

UzABSA-LLM: Parameter-Efficient Fine-Tuning and Multi-Domain Annotation for Uzbek Aspect-Based Sentiment Analysis

Mersaid Aripov ¹, Sanatbek Matlatipov ¹, Jaloliddin Rajabov ¹ and Abdullah M. Almarashi²

¹ Faculty of Applied Mathematics and Intelligent Technologies, National University of Uzbekistan named after Mirzo Ulugbek, Tashkent 100174, Uzbekistan; s.matlatipov@nuu.uz

² Department of Statistics, Faculty of Science, King Abdulaziz University, Jeddah 21589, Saudi Arabia; aalmarashi@kau.edu.sa

* Correspondence: s.matlatipov@nuu.uz; aalmarashi@kau.edu.sa

Abstract

Uzbek, a Turkic language with almost 38 million speakers, remains under-served in aspect-based sentiment analysis (ABSA). We present, to our knowledge, the first systematic study of parameter-efficient fine-tuning of open-source large language models (LLMs) for Uzbek ABSA: Qwen 2.5-7B, Llama 3.1-8B, and DeepSeek-R1-Distill-7B are fine-tuned under identical QLoRA configurations to jointly extract aspect terms and classify sentiment as structured JSON. On a 609-example evaluation set, Qwen achieves the best LLM extraction F1 (0.660) with a 100% parse rate, while Llama leads in end-to-end pair F1 (0.581) and sentiment accuracy (0.886). Benchmarking against fine-tuned Uzbek BERT encoders reveals a task-dependent trade-off: a 110M-parameter encoder slightly surpasses the LLMs at span extraction, whereas the LLMs are clearly stronger at sentiment and joint prediction. We then propose a three-layer model-assisted annotation pipeline—local inference, LLM-as-Judge quality scoring, and quality-tiered assembly—and apply it to 5,038 business reviews from 23 domains, producing a multi-domain Uzbek ABSA dataset of 9,884 aspect–sentiment pairs. A native-speaker validation study calibrates the LLM judge against human ratings (Spearman $\rho = 0.66$ for overall quality) and contributes an 80-review gold-verified subset. All code, models, and datasets are publicly released.

Keywords: Aspect-Based Sentiment Analysis; Low-Resource NLP; Uzbek Language; Parameter-Efficient Fine-Tuning; QLoRA; Large Language Models; LLM-as-Judge; Multi-Domain Dataset

1. Introduction

Aspect-Based Sentiment Analysis (ABSA) identifies specific aspects and their polarities in text [1]. While high-resource languages dominate the field, Uzbek—a Turkic language with almost 38 million speakers¹—remains comparatively under-explored, owing to script duality (Latin/Cyrillic), agglutinative morphology, and limited representation in multilingual pre-training corpora. Recent 7–8B-parameter LLMs such as Qwen 2.5 [2], Llama 3.1 [3], and DeepSeek-R1 [4] offer cross-lingual transfer potential. Using QLoRA [5], we adapt these models on consumer hardware by reformulating ABSA as a single-pass JSON generation task.

¹ <https://www.worldometers.info/world-population/uzbekistan-population/>

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the authors.

Submitted to *Computers* for possible

open access publication under the

terms and conditions of the [Creative](#)

[Commons Attribution \(CC BY\)](#) license.

This work is guided by three questions. First, we ask to what extent parameter-efficient fine-tuning with QLoRA enables open-source LLMs to perform joint aspect extraction and sentiment classification for Uzbek. Second, we examine how the chosen architectures compare under identical QLoRA configurations, and whether validation loss is a reliable guide to downstream task performance. Third, we ask whether a model trained only on restaurant reviews can produce usable annotations across a wider range of business domains, and how quality varies with distance from the training domain. In addressing these questions, we make the following contributions:

1. To our knowledge, the first systematic evaluation of QLoRA-fine-tuned open-source LLMs for Uzbek ABSA, comparing three architectures at the same parameter scale under identical training conditions, and benchmarked against fine-tuned Uzbek BERT encoders.
2. A scalable model-assisted annotation pipeline that combines local inference, LLM-as-Judge quality scoring, and quality-tiered assembly, together with a native-speaker validation study (a two-stage annotation-agreement analysis, calibration of the LLM judge, and a gold-verified subset) that quantifies the reliability of the resulting resource.
3. A multi-domain Uzbek ABSA dataset covering 5,038 reviews across 23 business categories with 9,884 aspect-sentiment pairs and per-annotation quality metadata. The present experiments use the original five-category UzABSA label space; a separate 119-subcategory taxonomy is released as an auxiliary resource for future domain-adaptive annotation.

A recurring observation across our experiments is that training cross-entropy loss does not track ABSA task performance: the model with the lowest validation loss is not the strongest on aspect extraction. While unsurprising for structured-generation tasks, this has a practical consequence for model selection, which we return to in Section 6. To support reproducible research in low-resource NLP, all code, datasets, and fine-tuned models are publicly available on GitHub² and HuggingFace³.

The remainder of this paper is organized as follows. Section 2 reviews related work on ABSA, parameter-efficient fine-tuning for low-resource languages, and LLM-based evaluation. Section 3 describes the two datasets used in this study. Section 4 presents the methodology, including model selection, the QLoRA configuration, and the annotation pipeline. Section 5 reports the experimental setup and results. Section 6 discusses implications and limitations, and Section 7 concludes with directions for future work.

2. Related Work

SemEval-2014 Task 4 [1] established the modern ABSA paradigm by defining Aspect Term Extraction (ATE) [6], Aspect Category Detection (ACD) [7], Aspect Sentiment Classification (ASC) [8], and joint end-to-end ABSA on English restaurant and laptop reviews [9]. Early systems typically combined sequence labeling (e.g., CRF/BiLSTM) for extraction with attention-based sentiment classifiers. With the rise of pre-trained Transformers, BERT-style models quickly dominated ABSA benchmarks [10,11], casting ATE as token classification and ASC as sentence(-pair) classification via fine-tuning.

Recent work increasingly reframes ABSA as generation: Scaria et al. [12] propose InstructABSA, using instruction-tuned LLMs to jointly extract aspect terms and predict sentiment in a single step by generating structured text, enabling unified modeling and more open-ended aspect vocabularies. However, generative ABSA has largely focused on

² <https://github.com/SanatbekMatlatipov/UzABSA-LLM>

³ <https://huggingface.co/Sanatbek/UzABSA-LLM>

English, leaving morphologically rich, low-resource languages underexplored; we extend this paradigm to Uzbek by fine-tuning open-source instruction-tuned LLMs to produce structured JSON outputs.

LLM Fine-Tuning for Low-Resource Languages. Large multilingual foundation models have broadened NLP beyond English. Architectures such as mT5, BLOOM, Llama [3,13], and Qwen [2] differ in low-resource coverage, largely reflecting their pre-training data; for Turkic languages, Qwen 2.5 reports broader multilingual coverage than Llama 3.1, but neither explicitly guarantees Uzbek support. Task adaptation in low-resource settings is constrained by the cost of full fine-tuning for 7–8B models (tens of GB VRAM and substantial compute). Parameter-efficient methods mitigate this: LoRA [14] adds trainable low-rank adapters to attention while freezing base weights, and QLoRA [5] combines LoRA with 4-bit NormalFloat quantization to reduce memory to < 6 GB for 7B-scale models, enabling single-GPU fine-tuning. Prior work reports strong results for parameter-efficient adaptation on several non-English/low-resource tasks (e.g., Arabic sentiment analysis, Turkish NER, Korean QA), and recent studies pair LoRA-tuned 7B models with encoder baselines for structured extraction, finding the generative approach competitive with supervised BERT-style taggers [15]. Yet, to our knowledge, LoRA/QLoRA has not been applied to Uzbek NLP. We address this gap by systematically evaluating QLoRA-adapted LLMs for Uzbek ABSA under controlled conditions, benchmarked against fine-tuned monolingual Uzbek encoders.

Uzbek Natural Language Processing. Early work addressed core infrastructure: Aripov et al. [16] studied multilingual MT with Uzbek, highlighting challenges from agglutinative morphology, while [17,18] introduced UzbekTagger, a rule-based POS tagger and one of the few public morphological tools. Script variation is another persistent obstacle: after the 1993 shift from Cyrillic to Latin, much user text still appears in Cyrillic; Salaev et al.[19] proposed a transliteration tool to normalize across scripts.

For sentiment analysis, Kuriyozov and Matlatipov [20] created an early sentence-level Uzbek dataset with baseline classifiers, later extended by [21] through a broader study of low-resource sentiment dataset construction and benchmarking. Matlatipov et al.[22,23] further examined restaurant-domain sentiment classification on locally collected reviews. The closest resource to our study is UzABSA [24], the first Uzbek ABSA dataset, following the SemEval-2014 Task 4 scheme [1], with 6,175 annotated restaurant sentences, 7,412 aspect-term spans with polarity, and 7,724 aspect-category labels over five categories (*ovqat, muhit, xizmat, narx, boshqalar*).

On the modeling side, monolingual Uzbek encoders have recently become available, notably BERTbek [25], trained on Uzbek news text, and TahrirchiBERT [26], trained on ~5B tokens of Latin-script Uzbek; we employ both as fine-tuned baselines in this study (Section 4.2). Overall, prior Uzbek NLP work, including UzABSA [24], has mainly relied on rule-based methods, traditional machine learning, or encoder-based neural models; to the best of our knowledge, no published work has fine-tuned LLMs (full or parameter-efficient) for Uzbek ABSA, a gap directly addressed by the present study.

LLM-as-Judge for Annotation Quality. LLM-as-Judge uses large language models as automated evaluators, providing a scalable alternative to human assessment of generated text. Zheng et al.[27] introduced MT-Bench and Chatbot Arena, showing that strong LLMs (e.g., GPT-4) can yield scores that correlate with human preferences in open-ended dialogue. This approach has since been extended to instruction-following, summarisation, and translation, offering cost efficiency, reproducibility, and scalability when human annotation is expensive or impractical—particularly for low-resource languages. However, LLM judges may suffer from positional bias, self-enhancement bias, and session-to-session

Table 1. Statistics of the annotated UzABSA dataset used for training and evaluation.

Dataset statistics		Term-level polarity distribution	
Original examples	6,175	Positive	4,153 (56.0%)
Usable ChatML examples	6,089	Neutral	1,601 (21.6%)
Aspect term annotations	7,412	Negative	1,555 (21.0%)
Aspect category annotations	7,724	Conflict	103 (1.4%)
Unique categories	5	Train / Val split	5,480 / 609
Avg. words per sentence	6.42		

score drift [28]. While these issues are documented for chat-style evaluation, they are less studied for structured tasks such as information extraction or ABSA.

We use LLM-as-Judge for a different goal: assessing structured ABSA annotations produced by a fine-tuned model in a low-resource setting without gold references. Concretely, GPT-4o-mini scores Uzbek aspect–sentiment outputs on five rubric dimensions (completeness, accuracy, sentiment correctness, relevance, overall quality) using a 1–5 Likert scale. Unlike prior work, our target is structured extraction rather than free-form text, and the judge must understand Uzbek input; we then apply stratified sampling over 23 domains and quality thresholds to guide dataset assembly (Section 4.3).

3. Datasets

This study draws on two complementary corpora. The first is the UzABSA dataset [24], a publicly available collection of 6,175 annotated Uzbek restaurant review sentences hosted on HuggingFace.⁴ After preprocessing and converting the original 6,175 examples into ChatML-formatted instruction–response pairs, we retain 6,089 usable training instances, which we split 90/10 into 5,480 training and 609 validation examples (seed=42). Each sentence is annotated following the SemEval-2014 Task 4 scheme [1]: aspect terms are marked with character-level spans and assigned one of four polarity labels (*positive*, *negative*, *neutral*, *conflict*), and aspect categories are independently labeled with polarity. The dataset contains 7,412 aspect term annotations and 7,724 aspect category annotations distributed across five domain-specific categories: *ovqat* (food), *muhit* (ambiance), *xizmat* (service), *narx* (price), and *boshqalar* (other). On average, each sentence contains 2.53 aspect annotations and 6.42 words.

A notable characteristic of the dataset is class imbalance: positive polarity accounts for 56% of term-level annotations (4,153 of 7,412), followed by neutral (21.6%), negative (21.0%), and conflict (1.4%). This skew toward positive sentiment reflects the natural distribution of restaurant reviews but may bias model predictions toward the majority class. We report macro-averaged sentiment F1 alongside micro-averaged extraction scores to partially account for this imbalance (Section 5). Table 1 summarizes the key statistics.

3.1. Multi-Domain Uzbek Review Corpus for Model-Assisted Annotation

The second corpus is a novel collection of business reviews scraped from Sharh,⁵ a publicly accessible Uzbek-language business review platform. After deduplication and minimum-length filtering (removing entries with fewer than five words), the corpus comprises 5,038 reviews spanning 630 unique businesses across 23 business domains. This corpus is not used for training; it serves exclusively as the annotation target for the model-assisted pipeline described in Section 4.3.

⁴ <https://huggingface.co/datasets/Sanatbek/aspect-based-sentiment-analysis-uzbek>
⁵ <https://sharh.commeta.uz/en>

Table 2. Top-10 domain distribution in the multi-domain Uzbek review corpus (5,038 reviews total).

Rank	Domain	Reviews	%
1	Restoran/Ovqatlanish (Restaurant/Dining)	1,626	32.3
2	Bank/Moliya (Banking/Finance)	577	11.5
3	Boshqa (Other)	507	10.1
4	To'lov tizimlari (Payment systems)	328	6.5
5	Elektron tijorat (E-commerce)	296	5.9
6	Ta'lim (Education)	287	5.7
7	Tibbiyot/Sog'liqni saqlash (Healthcare/Medicine)	285	5.7
8	Telekommunikatsiya (Telecommunications)	234	4.6
9	Oziq-ovqat do'konlari (Grocery stores)	215	4.3
10	Gul/Sovg'a (Flowers/Gifts)	173	3.4
	Remaining 13 domains	510	10.1
	Total	5,038	100.0

Data collection and usage were conducted with explicit permission from the data owner, as noted in Section 7. The annotated corpus derived from these reviews is publicly released on Zenodo to facilitate further research on low-resource ABSA.⁶ To support future domain-adaptive annotation work, we also release a domain-specific aspect sub-category taxonomy comprising 119 fine-grained sub-categories across 18 of the 23 business domains (e.g., *kredit*, *foizstavka*, *karta_xizmati* for banking; *shifokor_malakasi*, *diagnostika* for healthcare), available alongside the dataset on GitHub. This taxonomy is not used as the prediction label space in the current experiments; the fine-tuned models are evaluated using the original five UzABSA categories.

The domain distribution is highly skewed: *Restoran/Ovqatlanish* (restaurants and food service) is the largest domain with 1,626 reviews (32.3%), followed by *Bank/Moliya* (banking and finance) with 577 reviews (11.5%). The ten most frequent domains together account for approximately 89.9% of the corpus; Table 2 lists their counts and proportions.

Text statistics reflect the informal, often brief nature of online business reviews: the average review contains 13.55 words and 96.26 characters, considerably shorter than the restaurant review sentences in the UzABSA training corpus (6.42 words per sentence, which represent clause-level annotations rather than full reviews). A script-level analysis of the corpus reveals that 92.8% of reviews are written in Uzbek Latin script, 5.2% in Russian Cyrillic, and 2.0% in Uzbek Cyrillic—a distribution consistent with the ongoing post-Soviet script transition and the presence of a Russian-speaking minority among Uzbek internet users. No script normalization or transliteration was applied prior to annotation; the fine-tuned model was evaluated on its ability to handle this natural orthographic variation.

3.2. Task Formulation

We cast ABSA as conditional text generation rather than separate sequence labeling/classification. Given an Uzbek review, the model produces a single structured JSON containing all aspect terms with their categories and sentiment polarities, avoiding multi-stage pipelines, reducing error propagation, and enabling end-to-end optimization.

We use a ChatML-style instruction template with system/user/assistant roles: the system sets Uzbek-ABSA expertise, the user supplies the review and extraction directives, and the assistant returns the JSON annotation. The full formatting is shown in Figure 1.

The JSON output format serves multiple methodological purposes: (a) it ensures structured, parseable responses that enable automatic metric computation without post-

⁶ <https://doi.org/10.5281/zenodo.18790639>




Figure 1. Complete ChatML instruction template for Uzbek ABSA. The example demonstrates joint extraction of aspect terms (*ovqatlari* “food”, *narxlari* “prices”), category assignment (*ovqat*, *narx*), and sentiment classification in a single generation step.

processing; (b) it naturally accommodates variable numbers of aspects per review; and (c) it facilitates batch annotation pipelines for large-scale corpus creation. Each aspect annotation contains three fields: *term* (the extracted span), *category* (one of the five UzABSA categories: *ovqat*, *muhit*, *xizmat*, *narx*, *boshqalar*), and *polarity* (positive, negative, neutral, or conflict).

This generative formulation allows the model to leverage its pre-trained language understanding to simultaneously identify aspect boundaries, assign semantic categories, and classify sentiment polarity—tasks that traditional approaches handle sequentially. The instruction text is provided entirely in Uzbek to maintain language consistency and reduce code-switching effects that could degrade model performance on this low-resource language.

4. Methodology

Our design requires models with (i) open-source, permissively licensed weights, (ii) instruction-tuned variants for complex prompt following, (iii) ~7–8B parameters for controlled comparison under a fixed compute budget, and (iv) 4-bit quantized availability compatible with QLoRA. Under these constraints, we select three complementary LLMs: Qwen 2.5-7B-Instruct [2], Llama 3.1-8B-Instruct [3], and DeepSeek-R1-Distill-Qwen-7B (a Qwen-7B backbone distilled from DeepSeek-R1 [4]).

Qwen 2.5-7B-Instruct [2] is a 7.6B decoder-only transformer with multilingual pre-training (reported >29 languages) and a GQA+SwiGLU+RoPE configuration (28 heads, 4 KV heads), which may benefit Uzbek. Llama 3.1-8B-Instruct [3] is an 8B GQA+SwiGLU+RoPE model (32 heads, 8 KV heads) with different depth/width (32 layers; 4,096 hidden vs. Qwen’s 28 layers; 3,584), and its Turkic coverage is less explicitly documented. DeepSeek-R1-Distill-Qwen-7B matches Qwen architecturally but differs in weights due to reasoning-oriented distillation rather than standard instruction tuning [4]. This choice supports two controlled contrasts: Qwen vs. Llama isolates architectural effects under instruction tuning, while Qwen vs. Distill-Qwen isolates the pre-training objective

on a fixed architecture. All models are loaded from HuggingFace Hub in pre-quantized 4-bit NormalFloat via Unsloth,⁷ avoiding quantization as a confound.

4.1. QLoRA Fine-Tuning

We adopt QLoRA [5] for parameter-efficient adaptation. Base weights are quantized to 4-bit NormalFloat (NF4) with double quantization and BFloat16 compute. LoRA adapters ($r=16$, $\alpha=32$, dropout 0.05, no bias) are applied to all attention and feed-forward linear projections, yielding 40.4M trainable parameters ($\sim 0.92\%$ of total for the 7B models). Training is performed via the SFTTrainer with Unsloth's optimized fused kernels and gradient checkpointing to minimize memory footprint. All three models are trained under strictly identical hyperparameters to isolate architectural differences: 8-bit AdamW optimizer, learning rate of 2×10^{-4} with a cosine schedule and 0.1 warmup ratio, 0.01 weight decay, and gradient clipping at 1.0. We train for 1,000 steps on the 5,480-example split using a maximum sequence length of 2,048 tokens and an effective batch size of 16 (per-device batch 4, 4 accumulation steps).

4.2. Encoder Baselines

To situate the generative formulation against classical ABSA architectures, we fine-tune two monolingual Uzbek encoders as token-classification baselines: BERTbek [25], a BERT-base model pre-trained on Uzbek news text, and TahrirchiBERT [26], a RoBERTa-base model pre-trained on approximately 5B tokens of Uzbek Latin-script text. Both have ~ 110 M parameters—roughly 1.5% of the size of the 7–8B LLMs.

Joint extraction and polarity classification is cast as BIO tagging with polarity-augmented labels ($0, B-/I-\{positive, negative, neutral, conflict\}$), so a single encoder predicts aspect spans and their sentiment simultaneously; decoded spans are converted to the same {term, polarity} structure used by the generative models, enabling identical scoring. Gold character spans are projected onto whitespace/punctuation word tokens; 98.2% of training and 98.5% of validation aspect terms could be localized this way (unlocalizable terms are dropped from encoder *training* only—evaluation for all systems uses the full reference set). Both encoders are trained on the identical 5,480-example split used for the LLMs (learning rate 3×10^{-5} , 4 epochs, batch size 16, maximum length 192, seed 42) and evaluated on the same 609-example set. Notably, the encoder baselines require no GPU-server infrastructure: each fine-tuning run completes in approximately 10 minutes on an Apple M2 laptop.

4.3. Multi-Domain Annotation Pipeline

To build a silver-standard, multi-domain Uzbek ABSA dataset, we design a three-layer pipeline that combines local inference with external quality validation. This addresses a common low-resource constraint: for the 23 out-of-domain business categories, no gold annotations exist and large-scale Uzbek expert labeling is prohibitively costly.

Layer 1: Model Inference. We select the best-performing model from the three-model comparison (Section 5.1) as the annotation engine, prioritising operational reliability: 100% JSON parse rate, highest ATE recall (maximising aspect coverage), and fastest inference throughput. Under these criteria, Qwen 2.5-7B is chosen (Section 5.1). We run greedy batch inference over all 5,038 reviews using the same instruction template and hyperparameters as evaluation, with checkpoint-based crash recovery (state saved every 100 reviews).

Layer 2: LLM-as-Judge Quality Scoring. Without references, we use GPT-4o-mini (via the OpenAI API) as an external judge. We draw a stratified sample of 307 reviews across all 23 domains to cover both high-frequency domains (e.g., restaurants, $n=50$) and low-frequency domains (e.g., insurance, $n=3$). For each sampled review, the judge

⁷ <https://github.com/unslothai/unsloth>

assigns 1–5 Likert scores for five rubric dimensions: completeness (coverage of expressed opinions), accuracy (terms present in the text), sentiment correctness (polarity labels), relevance (category–domain fit), and overall quality. The judge relies only on Uzbek text comprehension, consistent with the reference-free LLM-as-Judge paradigm in Section 2.

Layer 3: Quality-Tiered Assembly. We partition outputs into *Include* (overall ≥ 3.5), *Flag for review* ($2.5 \leq \text{overall} < 3.5$), and *Exclude* (overall < 2.5). The assembly script joins predictions with quality metadata (per-dimension scores, business category, source URL) and exports: (i) full JSON with provenance metadata, (ii) a filtered silver-standard subset (included + unjudged reviews), and (iii) JSONL for direct loading via HuggingFace datasets. Non-sampled reviews are released with an “unjudged” flag; the 63.5% inclusion rate on the judged subset provides a statistical basis for their expected quality. The resulting corpus is released on Zenodo (DOI: [10.5281/zenodo.18790639](https://doi.org/10.5281/zenodo.18790639)).

4.4. Evaluation Framework

We evaluate model performance along four complementary dimensions, each targeting a distinct facet of the ABSA generation task.

Aspect Term Extraction (ATE). We report micro-averaged precision, recall, and F1 under two matching criteria. *Exact match* requires the predicted term to be identical to the gold term after case-folding and whitespace normalisation. *Partial match* relaxes this constraint: a prediction is counted as correct if the predicted string is a substring of the gold term or vice versa (i.e., a superstring relationship), accommodating morphological variation in Uzbek—for example, matching *ovqatlari* against *ovqat*. In both cases, metrics are aggregated across all test examples (micro-averaging) to avoid bias from examples with few or many aspect terms.

End-to-End ABSA (E2E). The primary evaluation metric is the *aspect–polarity pair F1*, which requires both the aspect term and its sentiment polarity to be correct for a true positive. This metric captures joint extraction and classification performance in a single score and is directly comparable across models.

Sentiment Classification. For aspect terms that are correctly matched (exact match), we additionally report sentiment accuracy and macro-averaged F1 over the polarity classes. This isolates the model’s sentiment assignment capability from its term extraction performance.

Instruction Following. We measure the *JSON parse success rate*—the percentage of model outputs that yield a syntactically valid JSON object containing the expected *aspects* key. Failed parses are excluded from ABSA metric computation; a low parse rate therefore both reduces the effective evaluation set and signals poor instruction compliance.

All inference is performed with near-greedy decoding (temperature=0.1, do_sample=False, max_new_tokens=512) to ensure deterministic, reproducible outputs. For the multi-domain annotation quality assessment, we supplement automatic metrics with the five-dimension LLM-as-Judge rubric described in Section 4.3.

5. Results

All experiments run on a single NVIDIA RTX A6000 GPU (48 GB VRAM) using Python 3.10 with Unsloth, TRL, PEFT, and bitsandbytes. The UzABSA dataset is split 90/10 into training (5,480) and validation (609) sets with a fixed seed of 42. Each of the three models is trained for exactly 1,000 gradient steps under the QLoRA configuration described in Section 4.1; the *only* variable across runs is the base model. We use the term validation set because the processed split was originally prepared for controlled model comparison; no additional hyperparameter search was performed on this split after observing downstream ABSA metrics. A fully separate held-out test split and zero-/few-shot prompting baselines

Table 3. Training efficiency on a single RTX A6000. All runs use identical QLoRA settings (40.4 M trainable parameters, 0.92% of total).

Metric	Qwen 2.5-7B	Llama 3.1-8B	DeepSeek-R1-7B
Compute time (min)	49	60	62
Throughput (samples/s)	5.44	4.46	4.28
GPU mem. allocated (GB)	5.4	5.6	5.5
Total FLOPs	1.24×10^{17}	1.50×10^{17}	1.42×10^{17}
Best eval loss	0.3656	0.2785	0.3363

Table 4. ABSA evaluation on the 609-example evaluation set. Best scores per metric are **bold**. LLMs use greedy decoding (temperature=0.1); encoder baselines (BERTbek, TahrirchiBERT) use BIO-polarity token classification (Section 4.2) and produce structured output by construction, so the JSON parse rate does not apply. Several inter-system gaps are not statistically significant (see text).

Task	Metric	Fine-tuned LLMs (QLoRA, 7–8B)			Encoders (~110M)	
		Qwen 2.5	Llama 3.1	DeepSeek-R1	BERTbek	Tahrirchi
ATE (exact)	Precision	0.6138	0.6394	0.5348	0.6595	0.6510
	Recall	0.7143	0.6713	0.6921	0.6796	0.6546
	F1	0.6603	0.6549	0.6034	0.6694	0.6528
ATE (partial)	Precision	0.7163	0.7411	0.6452	0.7833	0.7724
	Recall	0.8336	0.7781	0.8350	0.8072	0.7767
	F1	0.7705	0.7591	0.7279	0.7951	0.7746
E2E Pair	Precision	0.5387	0.5667	0.4448	0.5532	0.5490
	Recall	0.6269	0.5950	0.5756	0.5700	0.5520
	F1	0.5795	0.5805	0.5018	0.5615	0.5505
Sentiment	Accuracy	0.8777	0.8864	0.8317	0.8388	0.8432
	Macro-F1	0.8113	0.8435	0.7717	0.7691	0.7673
JSON parse	Success rate	100.0%	95.9%	95.4%	—	—

are left for future work; fine-tuned encoder baselines are evaluated on the same split (Section 4.2). Table 3 summarises the resulting training efficiency.

5.1. Three-Model ABSA Comparison

Table 4 presents the full evaluation on the 609-example evaluation set, covering the three QLoRA-fine-tuned LLMs and the two fine-tuned Uzbek encoder baselines (Section 4.2). Four complementary findings emerge.

Qwen excels at aspect extraction among the LLMs. Qwen 2.5-7B achieves the highest LLM ATE F1 on both exact (0.6603) and partial match (0.7705), driven by the best recall across all extraction metrics. It also attains a perfect 100% JSON parse rate—a critical advantage for production annotation pipelines.

Llama leads on sentiment and end-to-end performance. Llama 3.1-8B obtains the highest E2E Pair F1 (0.5805), sentiment accuracy (0.8864), and sentiment macro-F1 (0.8435) of any system. Its precision-oriented extraction compensates for slightly lower recall, yielding the best aspect–polarity pair matching. Notably, the Pair F1 gap between Llama and Qwen is only 0.001—effectively a tie.

DeepSeek underperforms despite reasoning distillation. DeepSeek-R1-Distill-Qwen-7B scores lowest on every ABSA metric. It over-predicts aspects substantially (933 predictions vs. 721 gold; ratio 1.29), whereas Llama is closest to the gold count (757; ratio 1.05) and Qwen moderately over-predicts (839; ratio 1.16). The R1 reasoning distillation, while designed for chain-of-thought tasks, appears detrimental to direct structured JSON extraction.

Encoders are competitive at extraction; LLMs win on sentiment. The fine-tuned Uzbek encoders are strong extraction baselines: BERTbek attains the best exact (0.6694) and partial (0.7951) ATE F1 of any system, despite having $\sim 1.5\%$ of the LLMs' parameters, and its exact-ATE margin over Qwen is small (+0.009). The LLM advantage concentrates elsewhere: sentiment macro-F1 (Llama 0.8435, Qwen 0.8113 vs. ≈ 0.77 for both encoders) and end-to-end pair F1 (0.5805/0.5795 vs. 0.5615/0.5505). A paired bootstrap (10,000 resamples) between the two encoders finds only their partial-match ATE difference significant ($p = 0.023$); their exact-ATE, pair-F1, and sentiment differences are statistical ties. These results indicate that for pure span extraction in Uzbek, a laptop-trainable 110M encoder remains a highly competitive choice, whereas the generative LLMs deliver their value in joint aspect-sentiment prediction, categorical labeling within the same output structure, and the instruction-following flexibility that the annotation pipeline in Section 4.3 depends on.

Training Loss vs. Task Performance. Consistent with our second research question, ranking the models by validation loss, as visualized in Fig. 2, diverges from ranking them by downstream task metrics. Llama achieves the lowest evaluation loss (0.2785) and the highest E2E Pair F1 (0.5805), yet Qwen surpasses it on exact ATE F1 (0.6603 vs. 0.6549) despite having a 31% higher validation loss (0.3656). Furthermore, DeepSeek holds the second-best evaluation loss (0.3363) but yields the worst ABSA scores across the board (ATE F1 0.6034; Pair F1 0.5018). Cross-entropy loss captures token-level prediction quality but is blind to structured output correctness—aspect boundary precision, JSON formatting, and polarity alignment. This finding implies that model selection for structured extraction tasks should not rely on validation loss alone.

Table 5. LLM-as-Judge aggregate quality scores (307-review stratified sample, 1–5 scale).

Dimension	Mean	Interpretation
Accuracy	4.47	Extracted terms are genuinely present
Sentiment	4.21	Polarity labels largely correct
Relevance	4.06	Categories contextually appropriate
Overall	3.75	Majority suitable for inclusion
Completeness	3.32	Main opinions captured; some missed




Figure 2. Training dynamics of the three fine-tuned LLMs over 1,000 global steps. (a) Validation loss curves for Qwen 2.5-7B, Llama 3.1-8B, and DeepSeek-R1-Distill-Qwen-7B. (b) Learning-rate schedule used during fine-tuning. (c) Training loss curves for the three models.

5.2. Multi-Domain Quality Assessment

To evaluate cross-domain generalisation, we apply the annotation pipeline (Section 4.3) using the Qwen model—selected for its 100% parse rate and highest ATE recall—and score a stratified sample of 307 reviews via GPT-4o-mini as LLM-as-Judge. Table 5 reports the aggregate quality scores across five dimensions.

Quality-tier assignment using the overall score yields 63.5% *Include* (≥ 3.5 ; $n=195$), 28.0% *Flag* ($2.5\text{--}3.49$; $n=86$), and 8.5% *Exclude* (< 2.5 ; $n=26$). The high accuracy score (4.47) confirms that the model rarely hallucinates aspects; the lower completeness score (3.32) reflects a tendency to miss subsidiary or domain-specific opinions.

Domain-Level Quality Gradient. Table 6 presents per-domain overall scores for representative domains spanning the quality spectrum; complete scores for all 23 domains

Table 6. Per-domain LLM-as-Judge quality scores (overall, 1–5 scale). Domains grouped by proximity to the training distribution.

Group	Domain	N	Comp.	Acc.	Sent.	Rel.	Overall
In-domain	Restoran	50	3.68	4.76	4.56	4.66	4.24
Near	Sayohat/Turizm	5	4.00	4.80	4.60	4.40	4.40
	Oziq-ovqat	15	3.47	4.33	4.47	4.13	3.80
Out-of-domain	Bank/Moliya	30	3.40	4.47	4.17	3.90	3.70
	Telekom.	25	3.04	4.36	3.64	3.64	3.32
	Ta’lim	20	2.90	4.20	3.80	3.45	3.25
Distant	Sug’urta	3	2.33	3.67	2.67	3.33	2.67
	Investitsiya	3	2.33	4.00	3.67	3.00	2.67

are provided in Table A1 (Section A). A clear gradient emerges from in-domain to out-of-domain categories: the model trained exclusively on restaurant reviews achieves its highest quality on *Restoran* (4.24) and near-domain categories sharing overlapping aspect vocabularies (e.g., *Sayohat/Turizm*, 4.40), while performance degrades on domains with specialised terminology such as *Telekommunikatsiya* (3.32) and *Ta’lim* (3.25). The most distant domains—*Sug’urta* (insurance) and *Investitsiya* (investment)—drop to 2.67, driven by completeness failures rather than accuracy: the accuracy sub-score remains ≥ 3.67 across all 23 domains. Overall, 15 of 23 domains (65%) meet the inclusion threshold (≥ 3.5).

Annotation Pipeline Output. Batch annotation of the full 5,038-review corpus produces 9,884 aspects (average 1.96 per review) with a 100% JSON parse rate across all 23 domains. The polarity distribution is 75.1% positive, 16.5% negative, and 8.4% neutral—reflecting a natural skew in user-generated reviews. The aspect category distribution reveals a notable pattern: *boshqalar* (miscellaneous) dominates at 45.5%, followed by *ovqat* (17.9%), *xizmat* (15.9%), *muhit* (10.7%), and *narx* (9.0%).

5.3. Human Validation of Annotation Quality

Because both the silver annotations and their quality scores originate from models, we validate a subsample against native-speaker judgement. Two native Uzbek-speaking authors annotated a stratified subsample of 150 reviews drawn from the 307 judged reviews (all 23 domains represented, deliberately oversampling distant and low-scoring domains), performing two tasks: (i) rating the model’s predicted annotations on the same five-dimension 1–5 rubric used by the LLM judge, and (ii) re-annotating a subset of 80 reviews from scratch to obtain human gold references.

Gold annotation followed a two-stage protocol. In the first, fully independent pass, the two annotators’ span-level agreement was low (exact-match F1 0.277, partial 0.415)—not because they disagreed about opinions, but because their segmentation and normalization conventions diverged (lemmatized vs. surface forms; naming the referent rather than the mention); polarity agreement on matched terms was perfect. In the second stage the annotators reconciled conventions and produced a consolidated gold standard with high residual span agreement (exact F1 0.927, polarity 100%), released as the gold-verified subset alongside the silver corpus. This two-stage outcome is itself informative: even for native speakers, *what counts as the aspect span* is the dominant source of ABSA annotation disagreement in Uzbek, while polarity is essentially unambiguous.

Judge calibration. Table 7 compares GPT-4o-mini’s scores with the mean human rating on the same 150 reviews (rubric ratings were produced after joint calibration of the scale; the raters’ post-calibration consistency is high, with quadratic-weighted κ of 0.91–0.97 across dimensions). Rank correlations are moderate to strong on the dimensions that drive

Table 7. Calibration of the GPT-4o-mini judge against the mean rating of two native-speaker annotators on the same 150 reviews (Spearman ρ , all $p < 10^{-6}$; MAE on the 1–5 scale).

Dimension	ρ	MAE	Human mean	Judge mean
Sentiment	0.806	0.34	4.25	4.24
Relevance	0.667	0.50	3.85	3.97
Overall	0.657	0.55	3.84	3.75
Accuracy	0.433	0.54	4.19	4.44
Completeness	0.462	0.94	4.04	3.30

dataset assembly—sentiment ($\rho = 0.81$), relevance ($\rho = 0.67$), and the tier-defining overall score ($\rho = 0.66$; all $p < 10^{-6}$)—and the judge’s mean overall score (3.75) closely matches the human mean (3.84). Correlations are weaker for completeness and accuracy ($\rho \approx 0.43$ –0.46), where the judge is systematically harsher on completeness (mean 3.30 vs. 4.04). We conclude that the LLM judge is a usable, though imperfect, proxy for native-speaker quality judgement, and that its quality tiers rank annotations in an order humans broadly agree with.

Model vs. human gold. Against the consolidated gold references, the model achieves exact-match ATE F1 of 0.209–0.224 and partial-match F1 of 0.425–0.433 (per annotator), with 100% polarity agreement on exactly matched terms. Two reference points contextualize these scores. First, the model’s partial-match agreement (≈ 0.43) matches the level two humans reach *before* convention reconciliation (0.415), while falling well short of the reconciled ceiling (0.93)—the gap is concentrated in span segmentation, not in opinion detection or polarity. Second, across domain-proximity groups the human-measured quality reproduces the judge’s gradient, falling from 0.244 exact pair F1 on general out-of-domain reviews to 0.170 on the most distant domains. We note the scope of this study: the annotators are authors of this paper, and the sample sizes per domain are small; a larger validation with external annotators is left for future work.

6. Discussion

The Completeness Bottleneck. Completeness (mean 3.32/5) is the weakest quality dimension: the model reliably extracts aspects present in the text but often misses domain-specific or subtle opinions, especially outside the restaurant domain. This is primarily due to the training data’s narrow coverage—only five aspect categories—forcing the model to map unfamiliar aspects to the generic *boshqalar* (miscellaneous) label, which dominates at 45.5%. Addressing this bottleneck will require domain-adaptive fine-tuning or expanding the aspect taxonomy to better capture out-of-domain phenomena.

Architecture vs. Pre-Training Effects. DeepSeek’s R1 reasoning distillation does not benefit structured ABSA: despite sharing the Qwen-7B backbone, it underperforms both Qwen and Llama, likely due to reasoning-oriented output patterns conflicting with strict JSON formatting. Llama achieves the lowest loss and best sentiment accuracy, but Qwen is superior for reliable extraction and instruction following (100% parse rate, highest ATE recall). For annotation pipelines, robust JSON compliance and extraction coverage are more critical than marginal gains in sentiment metrics.

When is a 7B LLM Worth It? The encoder baselines sharpen the practical picture. A 110M monolingual encoder, fine-tuned in ten minutes on a laptop, matches or exceeds the 7–8B LLMs at aspect-term extraction—a reminder that for span identification in a morphologically rich language, targeted monolingual pre-training (BERTbek, TahrirchiBERT) remains remarkably effective. The LLMs justify their cost where the task is genuinely joint and generative: sentiment macro-F1 (+4–7 points over the encoders), unified term–category–polarity output, and the ability to follow natural-language instructions on unseen

domains, which is what the multi-domain annotation pipeline exploits. Teams building Uzbek ABSA systems under tight compute budgets should therefore consider encoders for extraction-only workloads and reserve LLMs for joint prediction and annotation-generation scenarios.

Loss as a Model Selection Proxy. Validation loss ranking diverges from ABSA performance: Llama’s lowest loss does not translate to the best extraction, while Qwen outperforms on ATE F1 despite higher loss. Cross-entropy loss captures token-level likelihood but not structured-output quality such as JSON validity or aspect-boundary precision. This divergence is expected for structured-generation tasks, yet it is easy to overlook in practice; model selection for ABSA is therefore better guided by task-specific metrics than by validation loss alone.

Limitations. This study trains exclusively on restaurant reviews, without domain-adaptive fine-tuning; cross-domain performance should therefore be read as a demonstration of pipeline feasibility rather than a claim of uniform quality, and indeed we observe measurable degradation on distant domains. While we validate a stratified subset of the multi-domain corpus against native-speaker annotations (Section 5.3), the released dataset remains predominantly silver-standard: the human-verified portion is an 80-review subset produced by two author–annotators under a two-stage reconciliation protocol, and the remaining annotations carry only model- and judge-derived quality estimates. Positive polarity dominates both datasets, and compute constraints meant the model comparison reuses the fixed evaluation split rather than a separately retrained held-out test set; for the same reason, statistical significance testing covers the encoder pair, while the LLM–encoder gaps are reported descriptively. Tokenizer efficiency and larger-scale, convention-reconciled human verification remain for future work.

7. Conclusions

This paper presented, to our knowledge, the first systematic study of parameter-efficient fine-tuning of open-source LLMs for Uzbek ABSA. We showed that QLoRA enables effective joint aspect extraction and sentiment classification via structured JSON generation, and that the fine-tuned LLMs are competitive with fine-tuned Uzbek BERT encoders. Our comparison revealed a clear architectural trade-off: Qwen 2.5-7B excels at reliable aspect extraction and instruction following, whereas Llama 3.1-8B achieves the highest precision for end-to-end sentiment matching, with the two effectively tied on end-to-end F1. We also found that validation loss is an unreliable proxy for downstream structured-extraction performance. Finally, we showed that a model trained only on restaurant reviews can produce usable annotations on more distant domains, though quality degrades measurably with domain distance. Our three-layer pipeline, together with a human-validation study of its output, produced a quality-scored multi-domain silver-standard resource with a gold-verified subset, offering a practical route to scalable low-resource dataset creation. Future work will focus on expanding the aspect taxonomy and on domain-adaptive fine-tuning to improve extraction completeness across diverse business domains.

Author Contributions: Conceptualization, M.A., S.M and A.M.; methodology, M.A and S.M.; software, S.M. and J.R.; validation, A.M.; investigation, M.A., S.M. and J.R.; resources, S.M.; data curation, J.R.; writing—original draft preparation, M.A. and S.M.; writing—review and editing, S.M. and A.M.; visualization, S.M.; supervision, M.A. and A.M.; project administration, M.A. and S.M; funding acquisition, not applicable. All authors have read and agreed to the published version of the manuscript.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable. The study analyses publicly accessible user-generated business reviews and does not involve direct interaction with human participants. The reviews were used with explicit permission from the platform owner (see Section 7).

Data Availability Statement: The primary Uzbek aspect-based sentiment analysis dataset used for model fine-tuning is publicly available on Hugging Face at <https://huggingface.co/datasets/Sanatbek/aspect-based-sentiment-analysis-uzbek>. The multi-domain Uzbek business review corpus generated in this study, including the 80-review gold-verified subset with independent double annotation (Section 5.3), is publicly available on Zenodo at <https://zenodo.org/records/18790639> with DOI: [10.5281/zenodo.18790639](https://doi.org/10.5281/zenodo.18790639). Code for the encoder baselines, the human-validation analysis, and the significance tests is available in the project repository.

Acknowledgments: We extend our heartfelt gratitude to Sardor Berdiyev, CEO and Co-founder of Commeta, for graciously granting permission to use publicly available reviews from the sharh.commeta.uz platform exclusively for academic research purposes. We also acknowledge the computational resources provided by the National University of Uzbekistan. “During the preparation of this manuscript/study, the author Sanatbek used OpenAI’s GPT model for the purposes of paraphrasing and better writing purposes. The authors have reviewed and edited the output and take full responsibility for the content of this publication.”

Conflicts of Interest: The authors declare no conflicts of interest.

Abbreviations

The following abbreviations are used in this manuscript:

ABSA	Aspect-Based Sentiment Analysis	
ACD	Aspect Category Detection	
ASC	Aspect Sentiment Classification	
ATE	Aspect Term Extraction	
QLoRA	Quantized Low-Rank Adaptation	
RQ	Research Question	

Appendix A. Full Per-Domain LLM-as-Judge Quality Scores

Table A1 reports LLM-as-Judge quality scores for all 23 business domains in the stratified evaluation sample (307 reviews total). The main text (Table 6) presents a selected subset grouped by training-distribution proximity; this table provides the complete picture, sorted by overall quality score in descending order. Scores are on a 1–5 Likert scale across five dimensions: Completeness, Accuracy (hallucination control), Sentiment Correctness, Relevance (category–domain fit), and Overall. Domains with $N \leq 5$ should be interpreted with caution due to small sample sizes. The “Include” threshold (Overall ≥ 3.5) is shown as a dashed reference.

Table A1. LLM-as-Judge (GPT-4o-mini) quality scores for all 23 business domains, sorted by Overall score (descending). Stratified evaluation sample: 307 reviews. Scale: 1–5 (higher is better). † Small sample ($N \leq 5$): interpret with caution.

Domain (Uzbek / English)	N	Comp.	Acc.	Sent.	Rel.	Overall
Bozor/BSC (Market/Shopping centre) [†]	3	3.67	5.00	5.00	5.00	4.67
Sayohat/Turizm (Tourism/Travel) [†]	5	4.00	4.80	4.60	4.40	4.40
Din/Madaniyat (Religion/Culture) [†]	3	4.33	5.00	4.00	5.00	4.33
Yetkazib berish (Delivery) [†]	3	3.67	5.00	4.67	4.33	4.33
Restoran/Ovqatlanish (Restaurant/Dining)	50	3.68	4.76	4.56	4.66	4.24
Go'zallik (Beauty services) [†]	3	3.33	4.67	5.00	4.00	4.00
Texnologiya/Media (Technology/Media)	8	2.88	4.50	5.00	3.88	3.88
Boshqa (Other/Miscellaneous)	37	3.43	4.54	3.89	4.38	3.81
Elektron tijorat (E-commerce)	21	3.43	4.48	4.10	4.14	3.81
Oziq-ovqat do'konlari (Grocery stores)	15	3.47	4.33	4.47	4.13	3.80
Tibbiyot/Sog'liqni saqlash (Healthcare)	25	3.40	4.60	4.20	4.36	3.80
Bank/Moliya (Banking/Finance)	30	3.40	4.47	4.17	3.90	3.70
Ko'ngilochar (Entertainment) [†]	3	3.33	4.67	4.33	4.00	3.67
Transport/Yo'l (Transport/Roads) [†]	3	3.33	4.00	4.00	4.00	3.67
Gul/Sovg'a (Flowers/Gifts)	12	3.08	4.08	4.83	3.67	3.58
To'lov tizimlari (Payment systems)	24	3.00	4.29	4.38	3.54	3.54
<i>Below inclusion threshold (Overall < 3.5)</i>						
Davlat xizmatlari (Government services) [†]	3	3.00	4.00	4.33	3.67	3.33
Telekommunikatsiya (Telecommunications)	25	3.04	4.36	3.64	3.64	3.32
Ta'lim (Education)	20	2.90	4.20	3.80	3.45	3.25
Sport/Fitnes (Gyms/Fitness) [†]	4	2.75	4.25	4.25	3.00	3.25
Kitob/Nashriyot (Books/Publishing) [†]	4	3.00	4.25	4.00	3.00	3.00
Sug'urta (Insurance) [†]	3	2.33	3.67	2.67	3.33	2.67
Investitsiya/Trading (Investment) [†]	3	2.33	4.00	3.67	3.00	2.67

References

- Pontiki, M.; Galanis, D.; Pavlopoulos, J.; Papageorgiou, H.; Androutsopoulos, I.; Manandhar, S. SemEval-2014 Task 4: Aspect Based Sentiment Analysis. In Proceedings of the Proceedings of the 8th International Workshop on Semantic Evaluation (SemEval 2014); Nakov, P.; Zesch, T., Eds., Dublin, Ireland, August 2014; pp. 27–35. <https://doi.org/10.3115/v1/S14-2004>.
- Qwen.; Yang, A.; Yang, B.; Zhang, B.; Hui, B.; Zheng, B.; Yu, B.; Li, C.; Liu, D.; Huang, F.; et al. Qwen2.5 Technical Report. *arXiv preprint* **2024**, [arXiv:cs.CL/2412.15115].
- Grattafiori, A.; Dubey, A.; Jauhri, A.; et al. The Llama 3 herd of models. *arXiv preprint* **2024**, [arXiv:cs.AI/2407.21783].
- DeepSeek-AI.; Guo, D.; Yang, D.; Zhang, H.; Song, J.; Wang, P.; et al. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *arXiv preprint* **2026**, [arXiv:cs.CL/2501.12948].
- Dettmers, T.; Pagnoni, A.; Holtzman, A.; Zettlemoyer, L. QLORA: efficient finetuning of quantized LLMs. In Proceedings of the Proceedings of the 37th International Conference on Neural Information Processing Systems, Red Hook, NY, USA, 2023; NIPS '23.
- Ma, D.; Li, S.; Wu, F.; Xie, X.; Wang, H. Exploring Sequence-to-Sequence Learning in Aspect Term Extraction. In Proceedings of the Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics; Korhonen, A.; Traum, D.; Màrquez, L., Eds., Florence, Italy, July 2019; pp. 3538–3547. <https://doi.org/10.18653/v1/P19-1344>.
- Hu, M.; Zhao, S.; Guo, H.; Xue, C.; Gao, H.; Gao, T.; Cheng, R.; Su, Z. Multi-Label Few-Shot Learning for Aspect Category Detection. In Proceedings of the Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers); Zong, C.; Xia, F.; Li, W.; Navigli, R., Eds., Online, August 2021; pp. 6330–6340. <https://doi.org/10.18653/v1/2021.acl-long.495>.
- Sun, X.; Mi, Y.; Li, H. Enhancing Aspect Sentiment Classification with Dual-Channel Graph Convolutional Network. *ACM Trans. Intell. Syst. Technol.* **2025**, *16*. <https://doi.org/10.1145/3721844>.
- Hua, Y.C.; Denny, P.; Wicker, J.; Taskova, K. A systematic review of aspect-based sentiment analysis: domains, methods, and trends. *Artificial Intelligence Review* **2024**, *57*, 296. <https://doi.org/10.1007/s10462-024-10906-z>.
- Zohreh Madhoushi, Z.M.; Hamdan, A.R.; Zainudin, S. Aspect-Based Sentiment Analysis Methods in Recent Years. *Asia-Pacific Journal of Information Technology; Multimedia* **2019**, *08*, 79–96. <https://doi.org/10.17576/apjitm-2019-0801-07>.
- Chouikhi, H.; Alsuhaibani, M.; Jarray, F. BERT-Based Joint Model for Aspect Term Extraction and Aspect Polarity Detection in Arabic Text. *Electronics* **2023**, *12*. <https://doi.org/10.3390/electronics12030515>.

12. Scaria, K.; Gupta, H.; Goyal, S.; Sawant, S.; Mishra, S.; Baral, C. InstructABSA: Instruction Learning for Aspect Based Sentiment Analysis. In Proceedings of the Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers); Duh, K.; Gomez, H.; Bethard, S., Eds., Mexico City, Mexico, June 2024; pp. 720–736. <https://doi.org/10.18653/v1/2024.naacl-short.63>.
13. Touvron, H.; Lavril, T.; Izacard, G.; Martinet, X.; Lachaux, M.A.; Lacroix, T.; Rozière, B.; Goyal, N.; Hambro, E.; Azhar, F.; et al. LLaMA: Open and efficient foundation language models. *arXiv preprint* **2023**, [[arXiv:cs.CL/2302.13971](https://arxiv.org/abs/2302.13971)].
14. Hu, E.J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; Chen, W. LoRA: Low-Rank Adaptation of Large Language Models. In Proceedings of the International Conference on Learning Representations, 2022.
15. Zhang, K.; Lin, X.; Wang, M.; Yu, H.; Chen, L. Named Entity Recognition Method for Natural Disaster Emergencies Based on Instruction Tuning and Graph Retrieval-Augmented Generation. *Big Data and Cognitive Computing* **2026**, *10*, 185. <https://doi.org/10.3390/bdcc10060185>.
16. Aripov, M.; Khakimov, M.; Matlatipov, S.; Sirojiddinov, Z. Analysis and Processing of the Uzbek Language on the Multi-language Modelled Computer Translator Technology. In Proceedings of the Human Language Technology. Challenges for Computer Science and Linguistics; Vetulani, Z.; Paroubek, P.; Kubis, M., Eds., Cham, 2022; pp. 81–95.
17. Sharipov, M.; Kuriyozov, E.; Yuldashev, O.; Sobirov, O. UzbekTagger: The rule-based POS tagger for Uzbek language, 2023, [[arXiv:2301.12711](https://arxiv.org/abs/2301.12711)].
18. Sharipov, M.; Kuriyozov, E.; Yuldashev, O.; Sobirov, O. UzbekTagger: The rule-based POS tagger for Uzbek language. In Proceedings of the 10th Language & Technology Conference: Human Language Technologies as a Challenge for Computer Science and Linguistics, 2023.
19. Salaev, U. UzMorphAnalyser: A morphological analysis model for the Uzbek language using inflectional endings. *AIP Conference Proceedings* **2024**, *3244*, 030058. <https://doi.org/10.1063/5.0241461>.
20. Kuriyozov, E.; Matlatipov, S. Building a new sentiment analysis dataset for Uzbek language and creating baseline models. In Proceedings of the Multidisciplinary Digital Publishing Institute Proceedings, 2019, Vol. 21, p. 37.
21. Kuriyozov, E.; Matlatipov, S.; Alonso, M.A.; Gómez-Rodríguez, C. Construction and Evaluation of Sentiment Datasets for Low-Resource Languages: The Case of Uzbek. *Lecture Notes in Artificial Intelligence* **2022**, *13212*, 232–243. https://doi.org/10.1007/978-3-031-05328-3_15.
22. Matlatipov, G.; Vetulani, Z. Representation of Uzbek Morphology in Prolog. In *Aspects of Natural Language Processing: Essays Dedicated to Leonard Bolc on the Occasion of His 75th Birthday*; Marciniak, M.; Mykowiecka, A., Eds.; Springer Berlin Heidelberg: Berlin, Heidelberg, 2009; pp. 83–110. https://doi.org/10.1007/978-3-642-04735-0_4.
23. Matlatipov, S.; Tukeyev, U.; Aripov, M. Towards the Uzbek Language Endings as a Language Resource. In Proceedings of the Advances in Computational Collective Intelligence; Hernes, M.; Wojtkiewicz, K.; Szczerbicki, E., Eds., Cham, 2020; pp. 729–740.
24. Matlatipov, S.G.; Rajabov, J.; Kuriyozov, E.; Aripov, M. UzABSA: Aspect-Based Sentiment Analysis for the Uzbek Language. In Proceedings of the Proceedings of the 3rd Annual Meeting of the Special Interest Group on Under-resourced Languages@ LREC-COLING 2024, 2024, pp. 394–403.
25. Kuriyozov, E.; Vilares, D.; Gómez-Rodríguez, C. BERTbek: A Pretrained Language Model for Uzbek. In Proceedings of the Proceedings of the 3rd Annual Meeting of the Special Interest Group on Under-resourced Languages @ LREC-COLING 2024, Torino, Italy, 2024; pp. 33–44.
26. Mamasaidov, M.; Shopulatov, A. TahrirchiBERT base: A RoBERTa-based Language Model Pre-trained on 5B Tokens of Uzbek Latin-script Text. Hugging Face model card, 2023. Available online: <https://huggingface.co/tahrirchi/tahrirchi-bert-base> (accessed on 8 July 2026).
27. Zheng, L.; Chiang, W.L.; Sheng, Y.; Zhuang, S.; Wu, Z.; Zhuang, Y.; Lin, Z.; Li, Z.; Li, D.; Xing, E.P.; et al. Judging LLM-as-a-judge with MT-bench and Chatbot Arena. In Proceedings of the Proceedings of the 37th International Conference on Neural Information Processing Systems, Red Hook, NY, USA, 2023; NIPS '23.
28. Li, Q.; Dou, S.; Shao, K.; Chen, C.; Hu, H. Evaluating Scoring Bias in LLM-as-a-Judge, 2025, [[arXiv:2506.22316](https://arxiv.org/abs/2506.22316)].

Disclaimer/Publisher's Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.