

UzABSA-LLM: Parameter-Efficient Fine-Tuning and Multi-Domain Annotation for Uzbek Aspect-Based Sentiment Analysis

Sanatbek Matlatipov ^{1,*} , Gusein Djalilov ²  and Mersaid Aripov ¹ 

¹ Faculty of Applied Mathematics and Intelligent Technologies, National University of Uzbekistan named after Mirzo Ulugbek, Tashkent 100174, Uzbekistan; s.matlatipov@nuu.uz (S.M.); mersaid.aripov@nuu.uz (M.A.)

² Independent Researcher, Tashkent, Uzbekistan; g.djalilov@gmail.com

* Correspondence: s.matlatipov@nuu.uz

Abstract

We present a controlled study of parameter-efficient fine-tuning of open-weight large language models (LLMs) for Uzbek aspect-based sentiment analysis (ABSA). Qwen 2.5-7B, Llama 3.1-8B, and DeepSeek-R1-Distill-7B are fine-tuned under identical QLoRA configurations to jointly extract aspect terms and assign their sentiment as structured JSON, and evaluated on a 609-example review-grouped, deduplicated split against zero-shot controls. Fine-tuning at least doubles exact extraction F1 and triples aspect–polarity pair F1: Qwen achieves the best extraction F1 (0.708) with a 100% parse rate, while Llama leads on pair F1 (0.651) and sentiment accuracy (0.925); paired-bootstrap tests detect no significant difference between the two. Benchmarking against fine-tuned 110M-parameter Uzbek BERT encoders reveals a task-dependent trade-off: no significant difference is detected at aspect extraction, whereas the LLMs are clearly stronger at sentiment and joint prediction. We then propose a three-layer model-assisted annotation pipeline—local LLM annotation, LLM-as-Judge quality scoring, and quality-tiered assembly—and apply it to 5,038 business reviews from 23 domains, producing a multi-domain Uzbek ABSA dataset with 9,884 aspect-level annotations. A native-speaker validation study calibrates the judge against human ratings (Spearman $\rho = 0.66$ overall) and contributes an 80-review reconciled, double-annotated subset. All code, models, and datasets are publicly released.

Keywords: Aspect-Based Sentiment Analysis; Big Data Text Analytics; Low-Resource NLP; Uzbek Language; Parameter-Efficient Fine-Tuning; QLoRA; Large Language Models; LLM-as-Judge; Multi-Domain Dataset

1. Introduction

Aspect-Based Sentiment Analysis (ABSA) identifies certain aspects and polarity in text [1]. While high-resource languages dominate the field, Uzbek which is a Turkic language family with nearly 38 million speakers¹, is comparatively under-explored, owing to script duality (Latin/Cyrillic), agglutinative morphology, and limited representation in multilingual pre-training corpora. Recent 7 to 8B-parameter LLMs, such as Qwen 2.5 [2], Llama 3.1 [3], and DeepSeek-R1 [4], provide cross-lingual transfer potential. We use QLoRA [5] to adapt these models on consumer hardware, reformulating ABSA as a single-pass JSON generating operation.

¹ <https://www.worldometers.info/world-population/uzbekistan-population/>

Received:

Revised:

Accepted:

Published:

Copyright: © 2026 by the authors.

Submitted to *Big Data Cogn. Comput.*

for possible open access publication

under the terms and conditions of the

[Creative Commons Attribution](#)

(CC BY) license.

This work is guided by three questions. First, we ask to what extent parameter-efficient fine-tuning with QLoRA enables open-weight LLMs to perform joint aspect extraction and sentiment classification for Uzbek. Second, we examine how the chosen architectures compare under identical QLoRA configurations, and whether validation loss is a reliable guide to downstream task performance. Third, we ask whether a model trained only on restaurant reviews can produce usable annotations across a wider range of business domains, and how quality varies with distance from the training domain.

Beyond the language-specific results, the study addresses a problem generic to big-data text analytics: turning a heterogeneous, continuously growing stream of user-generated reviews into a quality-controlled analytical resource when expert annotation cannot scale with the data. Review platforms accumulate text far faster than annotators can label it, and in a low-resource language there is no pool of annotators to draw on in the first place—so the binding constraint on analysing this data is annotation throughput, not storage or compute. Our answer couples a locally fine-tuned annotator with an automated cognitive evaluation layer—an LLM judge whose scores we calibrate against native-speaker judgement—and a tiered assembly step that makes annotation quality explicit, per instance, in the released data. The pipeline is model-agnostic and transfers to other low-resource languages and domains where the same annotation bottleneck applies.

In addressing these questions, we make the following contributions:

1. To our knowledge, the first systematic evaluation of QLoRA-fine-tuned open-weight LLMs for Uzbek ABSA, comparing three architectures at the same parameter scale under identical training conditions on a deduplicated, review-grouped split, benchmarked against fine-tuned Uzbek BERT encoders and against unadapted zero-shot controls, with paired-bootstrap significance tests for every headline comparison.
2. A scalable model-assisted annotation pipeline that combines local inference, LLM-as-Judge quality scoring, and quality-tiered assembly, together with a native-speaker validation study (a two-stage annotation-agreement analysis, calibration of the LLM judge, and a reconciled double-annotated subset) that quantifies the reliability of the resulting resource.
3. A multi-domain Uzbek ABSA dataset covering 5,038 reviews across 23 business categories with 9,884 aspect-sentiment pairs and per-annotation quality metadata. The present experiments use the original five-category UzABSA label space; a separate 119-subcategory taxonomy is released as an auxiliary resource for future domain-adaptive annotation.

A recurring observation across our experiments is that training cross-entropy loss does not track ABSA task performance: the model with the lowest validation loss is not the strongest on aspect extraction. While unsurprising for structured-generation tasks, this has a practical consequence for model selection, which we return to in Section 6. To support reproducible research in low-resource NLP, all code, datasets, and fine-tuned models are publicly available on GitHub² and HuggingFace³.

The remainder of this paper is organized as follows. Section 2 reviews related work on ABSA, parameter-efficient fine-tuning for low-resource languages, and LLM-based evaluation. Section 3 describes the two datasets used in this study. Section 4 presents the methodology, including model selection, the QLoRA configuration, and the annotation pipeline. Section 5 reports the experimental setup and results. Section 6 discusses implications and limitations, and Section 7 concludes with directions for future work.

² <https://github.com/SanatbekMatlatipov/UzABSA-LLM>

³ <https://huggingface.co/Sanatbek/UzABSA-LLM>

2. Related Work

SemEval-2014 Task 4 [1] established the modern ABSA paradigm by defining Aspect Term Extraction (ATE) [6], Aspect Category Detection (ACD) [7], Aspect Sentiment Classification (ASC) [8], and joint end-to-end ABSA on English restaurant and laptop reviews [9]. Early systems typically combined sequence labeling (e.g., CRF/BiLSTM) for extraction with attention-based sentiment classifiers. With the rise of pre-trained Transformers, BERT-style models quickly dominated ABSA benchmarks [10,11], casting ATE as token classification and ASC as sentence(-pair) classification via fine-tuning, and neural sentiment classification architectures continue to evolve beyond the standard Transformer encoder [12]. The field has since broadened along two axes relevant to this study: richer structured targets combined with parameter-efficient fine-tuning, where updating under 1% of parameters approaches full fine-tuning on aspect–category–opinion–sentiment extraction [13], and non-English settings, where annotation scarcity has motivated unsupervised or distantly supervised ABSA pipelines, e.g., for Arabic traffic-service reviews [14].

Recent work increasingly reframes ABSA as generation: Scaria et al. [15] propose InstructABSA, using instruction-tuned LLMs to jointly extract aspect terms and predict sentiment in a single step by generating structured text, enabling unified modeling and more open-ended aspect vocabularies. Simmering and Huoviala [16] similarly benchmark GPT-family models against fine-tuned encoders on English ABSA, finding fine-tuned smaller models competitive with much larger prompted ones—a tension our encoder baselines revisit in a low-resource setting. Fine-tuned LLaMA-family models have since reached state-of-the-art results across English ABSA subtasks [17], whereas a systematic zero-shot evaluation of nine LLMs on multilingual ABSA finds prompting alone still trails fine-tuned baselines [18]—motivating the fine-tuning-versus-zero-shot controls in this study—and reasoning-infused multilingual ABSA models trained on corpora far larger than SemEval-2014 have recently been released for six higher-resource languages [19]. However, generative ABSA has largely focused on English and other well-resourced languages, leaving morphologically rich, low-resource languages underexplored; we extend this paradigm to Uzbek by fine-tuning open-weight instruction-tuned LLMs to produce structured JSON outputs.

LLM Fine-Tuning for Low-Resource Languages. Large multilingual foundation models have broadened NLP beyond English. Architectures such as mT5, BLOOM, Llama [3,20], and Qwen [2] differ in low-resource coverage, largely reflecting their pre-training data; for Turkic languages, Qwen 2.5 reports broader multilingual coverage than Llama 3.1, but neither explicitly guarantees Uzbek support. Task adaptation in low-resource settings is constrained by the cost of full fine-tuning for 7–8B models (tens of GB VRAM and substantial compute). Parameter-efficient methods mitigate this: LoRA [21] adds trainable low-rank adapters to attention while freezing base weights, and QLoRA [5] combines LoRA with 4-bit NormalFloat quantization to reduce memory to < 6 GB for 7B-scale models, enabling single-GPU fine-tuning. Prior work reports strong results for parameter-efficient adaptation on several non-English/low-resource tasks (e.g., Arabic sentiment analysis, Turkish NER, Korean QA), and recent studies pair LoRA-tuned 7B models with encoder baselines for structured extraction, finding the generative approach competitive with supervised BERT-style taggers [22]. Domain specialisation studies comparing fine-tuning and retrieval augmentation over the same open-weight backbones we adopt (Llama 3.1-8B, Qwen 2.5-7B) further show that adaptation strategy, not only model choice, drives downstream quality [23], while work on low-resource offensive-language detection highlights the robustness pitfalls of transferring models across data distributions [24]—a concern our multi-domain evaluation addresses directly. Yet, to our knowledge, LoRA/QLoRA has not been applied to Uzbek NLP. We address this gap by systematically evaluating

QLoRA-adapted LLMs for Uzbek ABSA under controlled conditions, benchmarked against fine-tuned monolingual Uzbek encoders.

Uzbek Natural Language Processing. Early work addressed core infrastructure: Matlatipov and Vetulani [25] modeled Uzbek morphology in Prolog, and Uzbek inflectional endings were later packaged as a stand-alone language resource [26]; Aripov et al. [27] studied multilingual MT with Uzbek, highlighting challenges from agglutinative morphology, while Sharipov et al. [28] introduced UzbekTagger, a rule-based POS tagger, and UzMorphAnalyser [29] added morphological analysis over inflectional endings. Script variation is another persistent obstacle: after the 1993 shift from Cyrillic to Latin, much user text still appears in Cyrillic; Salaev et al. [30] provide a machine transliteration tool between the Uzbek alphabets.

Syntactic annotation has begun to catch up with this morphological groundwork. A Universal Dependencies treebank of literary and educational Uzbek text now supplies gold morphosyntactic and dependency annotation for the language [31], and parsing studies built on it compare graph-based against transition-based architectures over BERT-based encoders [32] and show that cross-treebank training improves robustness when any single Uzbek treebank remains small [33]. This matters for ABSA because aspect terms in an agglutinative language are frequently nominal heads carrying case and possessive suffixes, and their opinion targets are identified by dependency relations rather than by adjacency—which is precisely the signal that syntax-aware ABSA taggers exploit [34]. The models compared in this study operate on raw text and use none of this structure; the availability of these resources is what would make syntax-aware Uzbek ABSA practical, and we return to it as future work in Section 7.

For sentiment analysis, Kuriyozov and Matlatipov [35] created an early sentence-level Uzbek dataset with baseline classifiers, later extended by [36] through a broader study of low-resource sentiment dataset construction and benchmarking. Matlatipov et al. [37] further examined sentiment classification of locally collected Uzbek restaurant reviews. The closest resource to our study is UzABSA [38], the first Uzbek ABSA dataset, following the SemEval-2014 Task 4 scheme [1], with 6,175 annotated restaurant sentences, 7,412 aspect-term spans with polarity, and 7,724 aspect-category labels over five categories (*ovqat, muhit, xizmat, narx, boshqalar*).

On the modeling side, monolingual Uzbek encoders have recently become available, notably BERTbek [39], trained on Uzbek news text, and TahrirchiBERT [40], trained on ~5B tokens of Latin-script Uzbek; we employ both as fine-tuned baselines in this study (Section 4.2).

Concurrently with this work, Yusufu et al. [34] released LASQ, a low-resource aspect-based sentiment *quadruple* dataset covering Uzbek and Uyghur, paired with a grid-tagging model that injects part-of-speech and dependency knowledge to counter the lexical sparsity of agglutinative languages. LASQ and the present study are complementary rather than competing: LASQ contributes new target–aspect–opinion–sentiment annotations over Google Play reviews and a syntax-aware discriminative tagger, whereas we hold the annotation scheme fixed (the existing SemEval-style UzABSA benchmark) and ask what parameter-efficient adaptation of open-weight generative LLMs delivers on it, relative to monolingual encoders. Their result that syntactic structure helps agglutinative extraction is a natural next ingredient for the models compared here.

Overall, prior Uzbek NLP work, including UzABSA [38], has mainly relied on rule-based methods, traditional machine learning, or encoder-based neural models. To the best of our knowledge, this is the first controlled comparison of QLoRA-adapted open-weight 7–8B models on the UzABSA benchmark, benchmarked against monolingual Uzbek encoders and coupled with a multi-domain annotation case study.

LLM-as-Judge for Annotation Quality. LLM-as-Judge uses large language models as automated evaluators, providing a scalable alternative to human assessment of generated text. Zheng et al.[41] introduced MT-Bench and Chatbot Arena, showing that strong LLMs (e.g., GPT-4) can yield scores that correlate with human preferences in open-ended dialogue. This approach has since been extended to instruction-following, summarisation, and translation, offering cost efficiency, reproducibility, and scalability when human annotation is expensive or impractical—particularly for low-resource languages. However, LLM judges may suffer from positional bias, self-enhancement bias, and session-to-session score drift [42]. While these issues are documented for chat-style evaluation, they are less studied for structured tasks such as information extraction or ABSA.

We use LLM-as-Judge for a different goal: assessing structured ABSA annotations produced by a fine-tuned model in a low-resource setting without gold references. Concretely, GPT-4o-mini scores Uzbek aspect-sentiment outputs on five rubric dimensions (completeness, accuracy, sentiment correctness, relevance, overall quality) using a 1–5 Likert scale. Unlike prior work, our target is structured extraction rather than free-form text, and the judge must understand Uzbek input; we then apply stratified sampling over 23 domains and quality thresholds to guide dataset assembly (Section 4.3). Closest to our pipeline, Hellwig et al. [43] train lightweight ABSA models directly on LLM-generated annotations, establishing model-assisted annotation as a viable route for ABSA; our pipeline differs in adding an explicit judging layer whose scores are themselves calibrated against native-speaker ratings (Section 5.5).

3. Datasets

This study draws on two complementary corpora. The first is the UzABSA dataset [38], a publicly available collection of 6,175 annotated Uzbek restaurant review sentences hosted on HuggingFace.⁴ After preprocessing and converting the original 6,175 examples into ChatML-formatted instruction-response pairs, we retain 6,089 usable training instances, which we split 90/10 into 5,480 training and 609 evaluation examples (seed 42); the split groups identical normalised review texts on the same side of the partition, for reasons quantified in Section 5.3. Each sentence is annotated following the SemEval-2014 Task 4 scheme [1]: aspect terms are marked with character-level spans and assigned one of four polarity labels (*positive*, *negative*, *neutral*, *conflict*), and aspect categories are independently labeled with polarity. The dataset contains 7,412 aspect term annotations and 7,724 aspect category annotations distributed across five domain-specific categories: *ovqat* (food), *muhit* (ambiance), *xizmat* (service), *narx* (price), and *boshqalar* (other). On average, each sentence contains 2.53 aspect annotations and 6.42 words.

A notable characteristic of the dataset is class imbalance: positive polarity accounts for 56% of term-level annotations (4,153 of 7,412), followed by neutral (21.6%), negative (21.0%), and conflict (1.4%). This skew toward positive sentiment reflects the natural distribution of restaurant reviews but may bias model predictions toward the majority class. The *conflict* label is dropped during preprocessing (Section 3.2), leaving a three-class polarity space for all experiments. We report macro-averaged sentiment F1 alongside micro-averaged extraction scores to partially account for this imbalance (Section 5). Table 1 summarizes the key statistics.

3.1. Multi-Domain Uzbek Review Corpus for Model-Assisted Annotation

The second corpus is a novel collection of business reviews scraped from Sharh,⁵ a publicly accessible Uzbek-language business review platform. After deduplication and

⁴ <https://huggingface.co/datasets/Sanatbek/aspect-based-sentiment-analysis-uzbek>

⁵ <https://sharh.commeta.uz/en>

Table 1. Statistics of the annotated UzABSA dataset used for training and evaluation.

Dataset statistics		Term-level polarity distribution	
Original examples	6,175	Positive	4,153 (56.0%)
Usable ChatML examples	6,089	Neutral	1,601 (21.6%)
Aspect term annotations	7,412	Negative	1,555 (21.0%)
Aspect category annotations	7,724	Conflict	103 (1.4%)
Unique categories	5	Train / Val split	5,480 / 609
Avg. words per sentence	6.42		

Table 2. Top-10 domain distribution in the multi-domain Uzbek review corpus (5,038 reviews total).

Rank	Domain	Reviews	%
1	Restoran/Ovqatlanish (Restaurant/Dining)	1,626	32.3
2	Bank/Moliya (Banking/Finance)	577	11.5
3	Boshqa (Other)	507	10.1
4	To'lov tizimlari (Payment systems)	328	6.5
5	Elektron tijorat (E-commerce)	296	5.9
6	Ta'lim (Education)	287	5.7
7	Tibbiyot/Sog'liqni saqlash (Healthcare/Medicine)	285	5.7
8	Telekommunikatsiya (Telecommunications)	234	4.6
9	Oziq-ovqat do'konlari (Grocery stores)	215	4.3
10	Gul/Sovg'a (Flowers/Gifts)	173	3.4
	Remaining 13 domains	510	10.1
	Total	5,038	100.0

minimum-length filtering (removing entries with fewer than five words), the corpus comprises 5,038 reviews spanning 630 unique businesses across 23 business domains. This corpus is not used for training; it serves exclusively as the annotation target for the model-assisted pipeline described in Section 4.3.

Data collection and usage were conducted with explicit permission from the data owner, as noted in Section 7. The annotated corpus derived from these reviews is publicly released on Zenodo to facilitate further research on low-resource ABSA.⁶ To support future domain-adaptive annotation work, we also release a domain-specific aspect sub-category taxonomy comprising 119 fine-grained sub-categories across 18 of the 23 business domains (e.g., *kredit*, *foizstavka*, *karta_xizmati* for banking; *shifokor_malakasi*, *diagnostika* for healthcare), available alongside the dataset on GitHub. This taxonomy is not used as the prediction label space in the current experiments; the fine-tuned models are evaluated using the original five UzABSA categories.

The domain distribution is highly skewed: *Restoran/Ovqatlanish* (restaurants and food service) is the largest domain with 1,626 reviews (32.3%), followed by *Bank/Moliya* (banking and finance) with 577 reviews (11.5%). The ten most frequent domains together account for approximately 89.9% of the corpus; Table 2 lists their counts and proportions.

Text statistics reflect the informal, often brief nature of online business reviews: the average review contains 13.55 words and 96.26 characters. These are whole reviews, and are therefore roughly twice as long as the UzABSA training units (6.42 words), which are clause- or sentence-level annotation spans rather than complete reviews—a length and granularity mismatch the annotation model must bridge at inference time. A script-level analysis of the corpus reveals that 92.8% of reviews are written in Uzbek Latin script, 5.2% in Russian Cyrillic, and 2.0% in Uzbek Cyrillic—a distribution consistent with the

⁶ <https://doi.org/10.5281/zenodo.22146677>

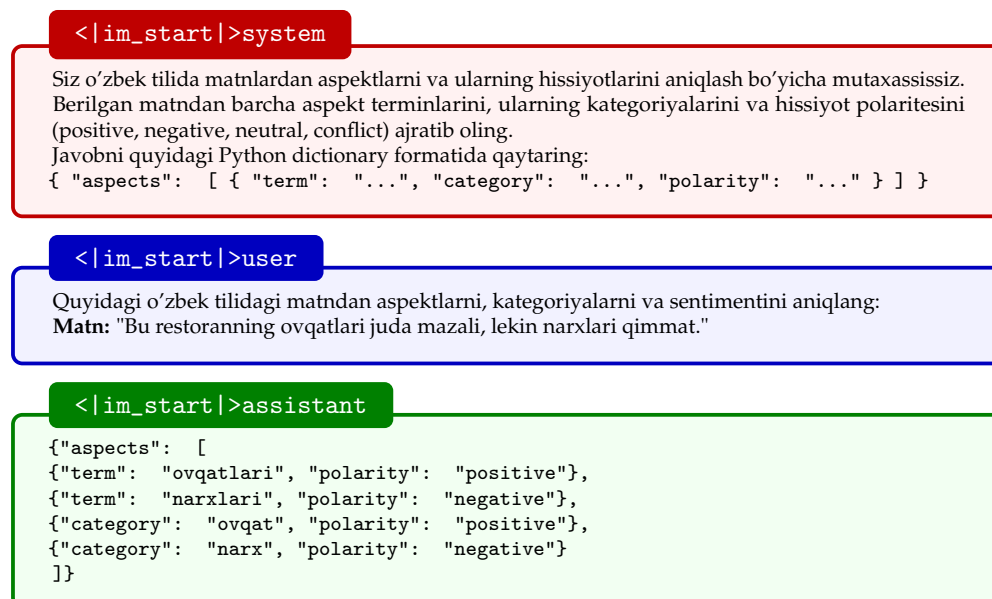


Figure 1. ChatML instruction template for Uzbek ABSA. The target is a single JSON list carrying *two* annotation layers, exactly as they appear in the SemEval-style source: term-level records (*ovqatlari* “food”, *narxlari* “prices”, each with a polarity) and category-level records (*ovqat*, *narx*, each with a polarity). The two layers are independently annotated in UzABSA and are therefore not paired into {term, category, polarity} triples; this study evaluates the term layer (Section 4.4).

ongoing post-Soviet script transition and the presence of a Russian-speaking minority among Uzbek internet users. No script normalization or transliteration was applied prior to annotation; the fine-tuned model was evaluated on its ability to handle this natural orthographic variation.

3.2. Task Formulation

We cast ABSA as conditional text generation rather than separate sequence labeling/classification. Given an Uzbek review, the model produces a single structured JSON containing all aspect terms with their categories and sentiment polarities, avoiding multi-stage pipelines, reducing error propagation, and enabling end-to-end optimization.

Formally, let x denote the instruction-formatted input (system prompt and review text) and let $y = (y_1, \dots, y_T)$ be the tokenized JSON annotation, where each aspect entry is a triple (t_j, c_j, p_j) of term, category, and polarity. Fine-tuning maximizes the conditional log-likelihood of the target sequence,

$$\mathcal{L}(\theta) = - \sum_{(x,y) \in \mathcal{D}} \sum_{t=1}^T \log p_{\theta}(y_t | y_{<t}, x), \quad (1)$$

where θ are the trainable adapter parameters (Section 4.1) and $\mathcal{D}$ is the instruction-formatted training set. At inference time the model generates $\hat{y}$ autoregressively with near-greedy decoding, and the aspect set $\{(\hat{t}_j, \hat{c}_j, \hat{p}_j)\}$ is recovered by JSON parsing.

We use a ChatML-style instruction template with system/user/assistant roles: the system sets Uzbek-ABSA expertise, the user supplies the review and extraction directives, and the assistant returns the JSON annotation. The full formatting is shown in Figure 1.

The JSON output format serves multiple methodological purposes: (a) it ensures structured, parseable responses that enable automatic metric computation without post-processing; (b) it naturally accommodates variable numbers of aspects per review; and (c) it facilitates batch annotation pipelines for large-scale corpus creation.

Two points about the target deserve emphasis, because they define what the reported metrics do and do not measure. First, UzABSA annotates aspect terms and aspect categories as *independent* layers: a category label is not attached to a particular term span. We preserve that independence rather than heuristically pairing the layers, so a target list contains term records (term, polarity) alongside category records (category, polarity), and the model learns to emit both. Consequently the “pair F1” reported throughout is an *aspect–polarity* pair, not a term–category–polarity triple, and category-layer accuracy is not scored here; constructing genuine triples would require a fresh annotation pass and is left for future work. Second, the four source polarity labels include *conflict*, which accounts for only 1.4% of term annotations (103 of 7,412). Our preprocessing retains only *positive*, *negative*, and *neutral*, so both training and evaluation—including the macro-averaged sentiment F1—operate over three classes.

This generative formulation allows the model to leverage its pre-trained language understanding to simultaneously identify aspect boundaries, assign semantic categories, and classify sentiment polarity—tasks that traditional approaches handle sequentially. The instruction text is provided entirely in Uzbek to maintain language consistency and reduce code-switching effects that could degrade model performance on this low-resource language.

4. Methodology

Our design requires models with (i) openly available (open-weight) checkpoints, (ii) instruction-tuned variants for complex prompt following, (iii) ~7–8B parameters for controlled comparison under a fixed compute budget, and (iv) 4-bit quantized availability compatible with QLoRA. Under these constraints, we select three complementary LLMs: Qwen 2.5-7B-Instruct [2], Llama 3.1-8B-Instruct [3], and DeepSeek-R1-Distill-Qwen-7B (a Qwen-7B backbone distilled from DeepSeek-R1 [4]).

Qwen 2.5-7B-Instruct [2] is a 7.6B decoder-only transformer with multilingual pre-training (reported >29 languages) and a GQA+SwiGLU+RoPE configuration (28 heads, 4 KV heads), which may benefit Uzbek. Llama 3.1-8B-Instruct [3] is an 8B GQA+SwiGLU+RoPE model (32 heads, 8 KV heads) with different depth/width (32 layers; 4,096 hidden vs. Qwen’s 28 layers; 3,584), and its Turkic coverage is less explicitly documented. DeepSeek-R1-Distill-Qwen-7B matches Qwen architecturally but differs in weights due to reasoning-oriented distillation rather than standard instruction tuning [4]. This choice supports two informative contrasts under a single fine-tuning protocol: Qwen vs. Llama, and Qwen vs. its R1-distilled sibling, the latter holding the architecture fixed. We stress that neither contrast isolates a single causal factor—the checkpoints also differ in tokenizer, pre-training corpus, instruction-tuning recipe, and (for Llama) parameter count—so differences are attributable to the checkpoints as wholes, not to architecture or distillation alone. All models are loaded from HuggingFace Hub in pre-quantized 4-bit NormalFloat via Unsloth,⁷ avoiding quantization as a confound.

4.1. QLoRA Fine-Tuning

We adopt QLoRA [5] for parameter-efficient adaptation. LoRA [21] freezes the pre-trained weights and injects a trainable low-rank decomposition into each adapted linear layer. For a frozen projection $W_0 \in \mathbb{R}^{d \times k}$, the adapted forward pass becomes

$$h = W_0x + \Delta Wx = W_0x + \frac{\alpha}{r} B A x, \quad A \in \mathbb{R}^{r \times k}, B \in \mathbb{R}^{d \times r}, \quad (2)$$

⁷ <https://github.com/unslothai/unsloth>

where $r \ll \min(d, k)$ is the adapter rank and α a scaling constant; only A and B are updated during training. QLoRA additionally stores W_0 in 4-bit NormalFloat (NF4) precision with double quantization, so the forward pass computes

$$h = \text{dequant}(W_0^{\text{NF4}})x + \frac{\alpha}{r}BAx \quad (3)$$

in BFloat16, reducing the memory footprint of a 7B-parameter model to under 6 GB while leaving gradient flow through the full-precision adapters intact [5].

In our configuration, adapters ($r=16$, $\alpha=32$, dropout 0.05, no bias) are applied to all attention and feed-forward linear projections. The resulting adapter size depends on the backbone's depth and width: 40.4M trainable parameters for the two Qwen-architecture models (28 layers, 3,584 hidden) and 41.9M for Llama 3.1-8B (32 layers, 4,096 hidden), corresponding to roughly 0.53% and 0.52% of the respective nominal parameter counts. Training is performed via the SFTTrainer with Unsloth's optimized fused kernels and gradient checkpointing to minimize memory footprint. All three models are trained under strictly identical hyperparameters: 8-bit AdamW optimizer, learning rate of 2×10^{-4} with a cosine schedule and 0.1 warmup ratio, 0.01 weight decay, and gradient clipping at 1.0. We train for 1,000 steps on the 5,480-example split using a maximum sequence length of 2,048 tokens and an effective batch size of 16 (per-device batch 4, 4 accumulation steps).

4.2. Encoder Baselines

To situate the generative formulation against classical ABSA architectures, we fine-tune two monolingual Uzbek encoders as token-classification baselines: BERTbek [39], a BERT-base model pre-trained on Uzbek news text, and TahrirchiBERT [40], a RoBERTa-base model pre-trained on approximately 5B tokens of Uzbek Latin-script text. Both have ~ 110 M parameters—roughly 1.5% of the size of the 7–8B LLMs.

Joint extraction and polarity classification is cast as BIO tagging with polarity-augmented labels ($0, B- / I-\{\text{positive}, \text{negative}, \text{neutral}, \text{conflict}\}$), so a single encoder predicts aspect spans and their sentiment simultaneously; decoded spans are converted to the same {term, polarity} structure used by the generative models, enabling identical scoring. Gold character spans are projected onto whitespace/punctuation word tokens; 98.2% of training and 98.5% of validation aspect terms could be localized this way (unlocalizable terms are dropped from encoder *training* only—evaluation for all systems uses the full reference set). Both encoders are trained on the identical 5,480-example split used for the LLMs (learning rate 3×10^{-5} , 4 epochs, batch size 16, maximum length 192, seed 42) and evaluated on the same 609-example set. Notably, the encoder baselines require no GPU-server infrastructure: each fine-tuning run completes in approximately 10 minutes on an Apple M2 laptop.

4.3. Multi-Domain Annotation Pipeline

To build a silver-standard, multi-domain Uzbek ABSA dataset, we design a three-layer pipeline that combines local inference with external quality validation. This addresses a common low-resource constraint: for the 23 out-of-domain business categories, no gold annotations exist and large-scale Uzbek expert labeling is prohibitively costly. Figure 2 gives a schematic overview, and Algorithm 1 states the procedure precisely.

Layer 1: Model Inference. We select the best-performing model from the three-model comparison (Section 5.1) as the annotation engine, prioritising operational reliability: 100% JSON parse rate, highest ATE recall (maximising aspect coverage), and fastest inference




Figure 2. The three-layer model-assisted annotation pipeline. A locally fine-tuned LLM annotates the full multi-domain corpus; an external LLM judge scores a stratified sample on a five-dimension rubric; assembly partitions the corpus into quality tiers. A native-speaker validation study (dashed) calibrates the judge and contributes a reconciled double-annotated subset.

Algorithm 1 Three-layer model-assisted annotation with quality tiering

Require: reviews $\mathcal{R} = \{x_1, \dots, x_N\}$ with domain labels d_i ; fine-tuned annotator M ; judge J ; sample size m ; thresholds $\tau_{\text{inc}} = 3.5$, $\tau_{\text{exc}} = 2.5$

Ensure: quality-tiered corpus $\mathcal{C}$

- 1: **Layer 1 (annotation):** for each $x_i \in \mathcal{R}$: $\hat{y}_i \leftarrow \text{parse}(M(x_i))$ ▷ greedy decoding;
checkpoint every 100 reviews
 - 2: **Layer 2 (judging):** draw stratified sample $\mathcal{S} \subset \mathcal{R}$, $|\mathcal{S}| = m$, covering every domain d
 - 3: **for** $x_i \in \mathcal{S}$ **do**
 - 4: $(s_i^{\text{comp}}, s_i^{\text{acc}}, s_i^{\text{sent}}, s_i^{\text{rel}}, s_i^{\text{ov}}) \leftarrow J(x_i, \hat{y}_i)$ ▷ 1–5 Likert per dimension
 - 5: **end for**
 - 6: **Layer 3 (assembly):** for each $x_i \in \mathcal{R}$:

$$\text{tier}(x_i) \leftarrow \begin{cases} \text{Include} & x_i \in \mathcal{S} \wedge s_i^{\text{ov}} \geq \tau_{\text{inc}} \\ \text{Flag} & x_i \in \mathcal{S} \wedge \tau_{\text{exc}} \leq s_i^{\text{ov}} < \tau_{\text{inc}} \\ \text{Exclude} & x_i \in \mathcal{S} \wedge s_i^{\text{ov}} < \tau_{\text{exc}} \\ \text{Unjudged} & x_i \notin \mathcal{S} \end{cases}$$
 - 8: **return** $\mathcal{C} = \{(x_i, \hat{y}_i, d_i, \text{tier}(x_i), \text{scores if judged})\}$
-

throughput. Under these criteria, Qwen 2.5-7B is chosen (Section 5.1).⁸ We run greedy batch inference over all 5,038 reviews using the same instruction template and hyperparameters as evaluation, with checkpoint-based crash recovery (state saved every 100 reviews).

Layer 2: LLM-as-Judge Quality Scoring. Without references, we use GPT-4o-mini (via the OpenAI API) as an external judge. We draw a stratified sample of 307 reviews across all 23 domains to cover both high-frequency domains (e.g., restaurants, $n=50$) and low-frequency domains (e.g., insurance, $n=3$). For each sampled review, the judge assigns 1–5 Likert scores for five rubric dimensions: completeness (coverage of expressed opinions), accuracy (terms present in the text), sentiment correctness (polarity labels), relevance (category–domain fit), and overall quality. The judge relies only on Uzbek text comprehension, consistent with the reference-free LLM-as-Judge paradigm in Section 2. Judging used OpenAI’s gpt-4o-mini (accessed via API on 25 February 2026) with temperature 0.1, a 300-token response cap, and one call per review; all 307 responses parsed successfully, and the full rubric prompt is released with the code.

Layer 3: Quality-Tiered Assembly. We partition judged outputs into *Include* (overall ≥ 3.5), *Flag for review* ($2.5 \leq \text{overall} < 3.5$), and *Exclude* (overall < 2.5). The assembly script joins predictions with quality metadata (per-dimension judge scores and business category) and exports: (i) full JSON with provenance metadata, (ii) a silver-standard subset comprising the judge-included reviews together with the non-sampled remainder, and (iii) JSONL for direct loading via HuggingFace datasets. Individual quality tiers exist only for the 307-review judged sample; every non-sampled review carries an explicit “unjudged” flag rather than a tier. Because the stratified sample deliberately over-represents low-

⁸ The case study was executed with the Qwen fine-tune from an earlier revision of the preprocessing pipeline (before the split and prompt-serialization fixes of Section 5.3). Because annotation quality is assessed directly by the judge layer and the human validation study, the pipeline’s conclusions are unaffected; see Section 6.

frequency domains, its raw 63.5% inclusion rate is a sample-level figure; reweighting the per-domain inclusion rates by the corpus domain distribution yields an estimated corpus-wide inclusion rate of 67.3% (95% CI [62.1%, 72.6%]). The resulting corpus is released on Zenodo (DOI: [10.5281/zenodo.22146677](https://doi.org/10.5281/zenodo.22146677)).

4.4. Evaluation Framework

We evaluate model performance along four complementary dimensions, each targeting a distinct facet of the ABSA generation task. Throughout, let G_i and P_i denote the gold and predicted aspect-term sets for review i , over an evaluation set of n reviews.

Aspect Term Extraction (ATE). We report micro-averaged precision, recall, and F1 under two matching criteria. *Exact match* requires the predicted term to be identical to the gold term after case-folding and whitespace normalisation. *Partial match* relaxes this constraint: a prediction is counted as correct if the predicted string is a substring of the gold term or vice versa (i.e., a superstring relationship), accommodating morphological variation in Uzbek—for example, matching *ovqatlari* against *ovqat*. Formally, with matching predicate $\mu \in \{\mu_{\text{ex}}, \mu_{\text{par}}\}$,

$$\mu_{\text{ex}}(t, g) = \mathbb{1}[\text{norm}(t) = \text{norm}(g)], \quad (4)$$

$$\mu_{\text{par}}(t, g) = \mathbb{1}[\text{norm}(t) \sqsubseteq \text{norm}(g) \vee \text{norm}(g) \sqsubseteq \text{norm}(t)], \quad (5)$$

where $\sqsubseteq$ denotes the substring relation, the micro-averaged scores are

$$P_\mu = \frac{\sum_{i=1}^n |\{t \in P_i : \exists g \in G_i, \mu(t, g)\}|}{\sum_{i=1}^n |P_i|}, \quad R_\mu = \frac{\sum_{i=1}^n |\{g \in G_i : \exists t \in P_i, \mu(t, g)\}|}{\sum_{i=1}^n |G_i|}, \quad (6)$$

and $F1_\mu = 2P_\mu R_\mu / (P_\mu + R_\mu)$. Micro-averaging across all test examples avoids bias from examples with few or many aspect terms. Two implementation properties should be noted: terms are normalized (lower-cased, whitespace-trimmed) and compared as sets, so repeated identical terms within one review are collapsed; and the existential predicates above do not enforce a one-to-one prediction–gold alignment for partial matches, so an optimal bipartite matching over character offsets would be a stricter variant, which we leave for future work. Both properties apply identically to every system evaluated, so cross-system comparisons are unaffected.

End-to-End ABSA (E2E). The primary evaluation metric is the *aspect–polarity pair F1*, which requires both the aspect term and its sentiment polarity to be correct for a true positive: Equation (6) is applied to the tuple sets $\{(t, p) : t \in P_i\}$ and $\{(g, p^*) : g \in G_i\}$ with the joint predicate $\mu_{\text{ex}}(t, g) \wedge \mathbb{1}[p = p^*]$. This metric captures joint extraction and classification performance in a single score and is directly comparable across models.

Sentiment Classification. For aspect terms that are correctly matched (exact match), we additionally report sentiment accuracy and macro-averaged F1 over the $K=3$ polarity classes retained after preprocessing,

$$\text{Acc} = \frac{|\{\text{matched terms with } p = p^*\}|}{|\{\text{matched terms}\}|}, \quad \text{macro-F1} = \frac{1}{K} \sum_{k=1}^K F1_k. \quad (7)$$

Because the polarity distribution is skewed (Section 3), macro-F1 weights the rarer *negative* and *neutral* classes equally with the dominant *positive* class (*conflict* is excluded in preprocessing, Section 3.2). This isolates the model’s sentiment assignment capability from its term extraction performance.

Instruction Following. We measure the *JSON parse success rate*—the percentage of model outputs that yield a syntactically valid JSON object containing the expected aspects

key. When a generation fails to parse, the evaluator falls back to a regular-expression salvage pass and scores whatever aspects it recovers, which is frequently none; failed parses are therefore counted as (usually empty) predictions rather than discarded. This is the conservative choice—it charges the model for its formatting failures as false negatives—but it means the parse rate and the ABSA metrics are not independent, and the encoders, which emit structured output by construction, are not exposed to this penalty at all.

Statistical Significance. Where two systems A and B are compared on the same evaluation set, we assess whether the observed metric difference $\hat{\delta} = m(A) - m(B)$ is distinguishable from zero using a paired bootstrap over reviews [44]: for $b = 1, \dots, B$ resamples ($B = 10,000$), we draw n reviews with replacement, recompute the metric for both systems on the identical resample, and record δ_b . We report the 95% percentile confidence interval of $\{\delta_b\}$ and the two-sided p-value

$$p = 2 \min \left(\frac{1}{B} \sum_b \mathbb{1}[\delta_b \leq 0], \frac{1}{B} \sum_b \mathbb{1}[\delta_b \geq 0] \right). \quad (8)$$

Pairing by review preserves the per-example correlation between systems and therefore tests the difference rather than the two absolute scores. All p-values are reported unadjusted. As a guard against multiple-comparison effects, we also apply a Holm–Bonferroni correction over the family of 28 post-hoc system-comparison tests (seven system pairs $\times$ four metrics): only the Llama–DeepSeek pair-F1 difference remains significant after adjustment (raw $p=0.0016$; corrected $p=0.045$), together with the three fine-tuned-versus-zero-shot contrasts (all raw $p<0.001$), which survive any standard correction. Individual sub-0.05 results elsewhere in Section 5 should therefore be read as exploratory evidence rather than confirmed differences.

All inference is performed with greedy decoding (`do_sample=False`, `max_new_tokens=512`), which is deterministic and reproducible. (The code additionally sets `temperature=0.1`, but this has no effect when sampling is disabled.) For the multi-domain annotation quality assessment, we supplement automatic metrics with the five-dimension LLM-as-Judge rubric described in Section 4.3.

5. Results

All experiments run on NVIDIA RTX A6000 GPUs (48 GB VRAM) using Python 3.13 with Unsloth, TRL, PEFT, and bitsandbytes; each run occupies a single GPU. The UzABSA dataset is split 90/10 into training (5,480) and evaluation (609) sets with a fixed seed of 42. The split is *grouped by review text*: all annotation records sharing a normalised review string are assigned to the same side, so no review can appear on both sides of the partition. This matters because the corpus annotates aspect terms and aspect categories as separate records (Section 3.2) and because short formulaic reviews recur verbatim, so a naive record-level random split leaks reviews across the partition; Section 5.3 quantifies what that leakage does to the reported numbers. Each of the three models is trained for exactly 1,000 gradient steps under the QLoRA configuration described in Section 4.1; the *only* variable across runs is the base model. We use the term evaluation set rather than test set because the same partition is monitored during training; no hyperparameter search was performed on it after observing downstream ABSA metrics. The fine-tuned encoder baselines (Section 4.2) and the unadapted zero-shot baselines (Section 5.2) are evaluated on this identical partition, so all systems in Table 4 are directly comparable. Table 3 summarises the resulting training efficiency.

Table 3. Training characteristics on a single RTX A6000. All runs use identical QLoRA hyperparameters; the adapter size differs slightly with backbone geometry (40.4 M trainable parameters for the Qwen-architecture models, 41.9 M for Llama 3.1-8B). Wall-clock time is reported only for the one run that was uninterrupted end-to-end; the other two runs were resumed from optimizer checkpoints after a power interruption, so their wall-clock and throughput are not clean efficiency measurements and are omitted rather than estimated (*).

Metric	Qwen 2.5-7B	Llama 3.1-8B	DeepSeek-R1-7B
Compute time (min)	—*	—*	84
GPU mem. allocated (GB)	5.4	5.6	5.5
Total FLOPs	2.19×10^{17}	2.49×10^{17}	2.37×10^{17}
Best eval loss (step)	0.1726 (700)	0.1476 (1000)	0.1744 (1000)

Table 4. ABSA evaluation on the 609-example deduplicated evaluation set. Best scores per metric are **bold**. LLMs use greedy decoding; encoder baselines (BERTbek, TahrirchiBERT) use BIO-polarity token classification (Section 4.2) and produce structured output by construction, so the JSON parse rate does not apply. Per-system percentile-bootstrap 95% confidence intervals for every score in this table are reported in Table A2 (Section B); their half-widths range from ± 0.023 to ± 0.037 , so most inter-system gaps are not statistically significant (see text and Section 5.2).

Task	Metric	Fine-tuned LLMs (QLoRA, 7–8B)			Encoders (~110M)	
		Qwen 2.5	Llama 3.1	DeepSeek-R1	BERTbek	Tahrirchi
ATE (exact)	Precision	0.7141	0.7185	0.7143	0.6802	0.7287
	Recall	0.7015	0.6893	0.6445	0.7042	0.6744
	F1	0.7077	0.7036	0.6776	0.6920	0.7005
ATE (partial)	Precision	0.8094	0.8175	0.8301	0.7864	0.8182
	Recall	0.7951	0.7843	0.7490	0.8141	0.7571
	F1	0.8022	0.8006	0.7874	0.8000	0.7865
E2E Pair	Precision	0.6506	0.6648	0.6361	0.5976	0.6471
	Recall	0.6391	0.6377	0.5739	0.6187	0.5997
	F1	0.6448	0.6510	0.6034	0.6080	0.6225
Sentiment	Accuracy	0.9110	0.9252	0.8905	0.8786	0.8893
	Macro-F1	0.8622	0.8867	0.8314	0.8126	0.8140
JSON parse	Success rate	100.0%	97.4%	96.9%	—	—

5.1. Three-Model ABSA Comparison

Table 4 presents the full evaluation on the 609-example evaluation set, covering the three QLoRA-fine-tuned LLMs and the two fine-tuned Uzbek encoder baselines (Section 4.2). Four complementary findings emerge.

Qwen leads extraction and is the most reliable formatter. Qwen 2.5-7B posts the highest exact (0.7077) and partial (0.8022) ATE F1 of any system and a perfect 100% JSON parse rate—a critical advantage for production annotation pipelines. A replicate of the Qwen run differing only in random seed reaches 0.7101/0.6554 (exact-ATE/pair F1 versus 0.7077/0.6448), so seed-level variation of roughly ± 0.01 should be kept in mind when reading small gaps in this table.

Llama leads on sentiment and end-to-end performance—but no significant difference from Qwen is detected. Llama 3.1-8B obtains the highest E2E pair F1 (0.6510), sentiment accuracy (0.9252), and sentiment macro-F1 (0.8867) of any system, with precision-oriented extraction compensating for slightly lower recall. However, a paired bootstrap over reviews ($B=10,000$; Equation (8)) finds no significant Qwen–Llama difference on any headline metric (exact-ATE $\Delta=+0.004$, $p=0.79$; pair F1 $\Delta=-0.006$, $p=0.68$; sentiment accuracy $\Delta=-0.014$, $p=0.27$): the observed gaps are within seed-level and sampling

variation—though absence of significance is not evidence of equivalence—and the choice between the two models can reasonably rest on secondary criteria such as parse reliability, licence, and tokenizer efficiency.

DeepSeek trails despite reasoning distillation. DeepSeek-R1-Distill-Qwen-7B scores lowest among the LLMs on every F1 metric (exact ATE 0.6776; pair F1 0.6034, below both Qwen, $p=0.009$, and Llama, $p=0.002$, under the paired bootstrap; the gap to Llama is the only pairwise difference that survives the Holm correction of Section 4.4). Its failure mode is recall: it finds the fewest gold aspects (0.6445 exact-ATE recall) while its precision matches the other two, and it under-predicts (665 predictions vs. 737 gold terms; ratio 0.90, versus 0.98 for Qwen and 0.96 for Llama). The R1 reasoning distillation, designed for chain-of-thought tasks, confers no advantage on direct structured extraction here—a pattern the zero-shot controls in Section 5.2 make considerably more vivid.

Encoders are competitive at extraction; LLMs win on sentiment. The fine-tuned Uzbek encoders remain strong extraction baselines at $\sim 1.5\%$ of the LLMs' parameters: TahrirchiBERT's exact-ATE F1 (0.7005) sits within 0.007 of Qwen's, and the paired bootstrap finds no significant extraction difference between the best LLM and either encoder (Qwen vs. TahrirchiBERT exact-ATE $p=0.61$; vs. BERTbek $p=0.29$), nor between the two encoders on any metric (exact-ATE $p=0.46$, partial-ATE $p=0.13$, pair F1 $p=0.25$, sentiment $p=0.39$). The joint-prediction advantage, by contrast, is real but uneven: Qwen's pair-F1 and sentiment margins over BERTbek are significant ($p=0.020$ and $p=0.021$), though not over TahrirchiBERT ($p=0.13$ and $p=0.14$), while Llama's sentiment-accuracy margin holds even against the stronger encoder ($\Delta=+0.036$ over TahrirchiBERT, $p=0.014$). These sub-0.05 joint-prediction results are unadjusted and do not survive the Holm correction of Section 4.4, so we read them as consistent exploratory evidence of an LLM advantage rather than individually confirmed differences. The LLM advantage concentrates elsewhere: sentiment macro-F1 (Llama 0.8867, Qwen 0.8622 vs. ≈ 0.81 for both encoders) and end-to-end pair F1 (0.6510/0.6448 vs. 0.6080/0.6225). For pure span extraction in Uzbek, a laptop-trainable 110M encoder therefore remains a highly competitive choice, whereas the generative LLMs deliver their value in joint aspect-sentiment prediction, categorical labeling within the same output structure, and the instruction-following flexibility that the annotation pipeline in Section 4.3 depends on.

Training Loss vs. Task Performance. Consistent with our second research question, ranking the models by validation loss, as visualized in Fig. 3, diverges from ranking them by downstream task metrics. Llama achieves the lowest evaluation loss (0.1476) and the highest E2E pair F1, yet Qwen surpasses it on exact ATE F1 despite a 17% higher evaluation loss (0.1726). More strikingly, Qwen and DeepSeek finish with nearly indistinguishable evaluation losses (0.1726 vs. 0.1744, a 1% gap) while differing by 3 points of exact-ATE F1 and 4 points of pair F1. Cross-entropy loss captures token-level prediction quality but is blind to structured output correctness—aspect boundary precision, JSON formatting, and polarity alignment. Model selection for structured extraction tasks should therefore not rely on validation loss alone.

5.2. What Does Fine-Tuning Buy? Zero-Shot Controls

To separate what QLoRA contributes from what the instruction-tuned base models already know, we evaluate each *unadapted* base model on the identical 609-example split with the identical Uzbek prompt and decoding settings. Table 5 reports the comparison.

Three observations follow. First, the fine-tuning effect is large and uniform: exact-ATE F1 rises by 0.38–0.46 and pair F1 by 0.43–0.45 absolute for every architecture, so the headline numbers in Table 4 are overwhelmingly a product of adaptation rather than of latent Uzbek ABSA competence in the base models. Second, fine-tuning buys *task* competence, not




Figure 3. Training dynamics of the three QLoRA fine-tuning runs over 1,000 gradient steps on the deduplicated split, plotted from the released per-step training logs. (a) Training loss (log scale, logged every 10 steps). (b) Evaluation loss on the 609-example set (every 100 steps): Llama 3.1-8B reaches the lowest loss (0.148), while Qwen 2.5-7B (0.173, minimum at step 700) and DeepSeek-R1-7B (0.174) are nearly indistinguishable—despite clearly different downstream ABSA scores (Table 4). (c) The cosine learning-rate schedule with 10% warmup shared by all runs.

Table 5. Zero-shot (ZS) versus QLoRA fine-tuned (FT) performance, same split, prompt, and decoding. Ratio is total predicted aspect terms over gold terms (737); values above 1 indicate over-prediction. Fine-tuning at least doubles exact-ATE F1 and triples pair F1 for every architecture; all FT–ZS differences are significant at $p < 0.001$ (paired bootstrap), surviving the Holm correction of Section 4.4.

Model		ATE F1 (ex.)	ATE F1 (par.)	Pair F1	Sent. macro-F1	JSON parse	Ratio
Qwen 2.5-7B	ZS	0.3222	0.5593	0.2100	0.4719	99.8%	1.80
	FT	0.7077	0.8022	0.6448	0.8622	100.0%	0.98
Llama 3.1-8B	ZS	0.2843	0.4924	0.1999	0.6280	95.4%	1.67
	FT	0.7036	0.8006	0.6510	0.8867	97.4%	0.96
DeepSeek-R1-7B	ZS	0.2169	0.3182	0.1540	0.5683	46.3%	1.10
	FT	0.6776	0.7874	0.6034	0.8314	96.9%	0.90

format compliance: Qwen already emits parseable JSON for 99.8% of zero-shot outputs, but attaches it to nearly twice as many aspect terms as the gold standard (ratio 1.80) with poorly calibrated polarities (sentiment macro-F1 0.4719). Fine-tuning’s clearest behavioural signature is this calibration of *how much* to extract—the prediction ratio moves from 1.10–1.80 to 0.90–0.98. Third, DeepSeek is the exception that proves the format rule: zero-shot, its R1-style reasoning traces break JSON parsing on more than half the examples (46.3%), and fine-tuning is what suppresses the reasoning preamble (96.9%)—yet even with formatting repaired, it remains the weakest of the three, indicating that reasoning distillation neither helps unadapted structured extraction nor survives as an advantage after adaptation.

5.3. Effect of Deduplication-Aware Splitting

Short formulaic reviews (*mazali* “tasty”, *ajoyib* “excellent”) recur verbatim in the UzABSA corpus, and the corpus stores aspect terms and aspect categories as separate annotation records (Section 3.2). A naive random split over records therefore leaks: in an earlier revision of this study’s pipeline, 61 of 609 evaluation items (10.0%) had an exact-text twin in training. Table 6 quantifies what review-grouped deduplication changes, using the encoder baselines—whose training pipeline is otherwise untouched by the fix—to isolate the split effect.

The direction of the change is instructive: contrary to the usual assumption that train–test leakage inflates scores, every system scores *higher* once leakage is removed. The mechanism lies in the record structure. Under the naive split, the term-layer and category-layer records of the same review could land on opposite sides, so a model repeatedly saw an

Table 6. Effect of replacing the naive record-level random split with review-grouped, deduplication-aware splitting (identical sizes: 5,480/609). Encoders isolate the split effect; the LLM deltas additionally include a prompt-serialization fix in the same pipeline revision. Every system *improves* on the leak-free split.

System	ATE exact F1			Pair F1		
	naive	grouped	Δ	naive	grouped	Δ
Qwen 2.5-7B	0.6603	0.7077	+0.047	0.5795	0.6448	+0.065
Llama 3.1-8B	0.6549	0.7036	+0.049	0.5805	0.6510	+0.071
DeepSeek-R1-7B	0.6034	0.6776	+0.074	0.5018	0.6034	+0.102
BERTbek	0.6694	0.6920	+0.023	0.5615	0.6080	+0.047
TahrirchiBERT	0.6528	0.7005	+0.048	0.5505	0.6225	+0.072

Table 7. LLM-as-Judge aggregate quality scores (307-review stratified sample, 1–5 scale).

Dimension	Mean	Interpretation
Accuracy	4.47	Extracted terms are genuinely present
Sentiment	4.21	Polarity labels largely correct
Relevance	4.06	Categories contextually appropriate
Overall	3.75	Majority suitable for inclusion
Completeness	3.32	Main opinions captured; some missed

eval-identical text in training paired with a *partial* target, and learned to reproduce exactly that partial answer at evaluation time. Deduplication also changes inferential conclusions, not just absolute scores: on the contaminated split, the two encoders’ partial-ATE difference tested as significant ($p=0.023$), a finding that does not survive on the clean split ($p=0.13$). Leakage can thus manufacture spurious model rankings even while depressing every absolute number—a caution we believe generalizes to other low-resource benchmarks built from short user-generated text, where verbatim duplication is endemic.

5.4. Multi-Domain Quality Assessment

To evaluate cross-domain generalisation, we apply the annotation pipeline (Section 4.3) using the Qwen model—selected for its 100% parse rate and best extraction F1 among the LLMs—and score a stratified sample of 307 reviews via GPT-4o-mini as LLM-as-Judge. Table 7 reports the aggregate quality scores across five dimensions.

Quality-tier assignment using the overall score yields 63.5% *Include* (≥ 3.5 ; $n=195$), 28.0% *Flag* (2.5–3.49; $n=86$), and 8.5% *Exclude* (< 2.5 ; $n=26$). These proportions describe the stratified sample, which intentionally over-samples low-frequency domains; reweighted to the corpus domain distribution, the estimated corpus-wide inclusion rate is 67.3% (95% CI [62.1%, 72.6%]). The high accuracy score (4.47) confirms that the model rarely hallucinates aspects; the lower completeness score (3.32) reflects a tendency to miss subsidiary or domain-specific opinions.

Domain-Level Quality Gradient. Table 8 presents per-domain overall scores for representative domains spanning the quality spectrum; complete scores for all 23 domains are provided in Table A1 (Section A). A clear gradient emerges from in-domain to out-of-domain categories: the model trained exclusively on restaurant reviews achieves its highest quality on *Restoran* (4.24) and near-domain categories sharing overlapping aspect vocabularies (e.g., *Sayohat/Turizm*, 4.40), while performance degrades on domains with specialised terminology such as *Telekommunikatsiya* (3.32) and *Ta’lim* (3.25). The most distant domains—*Sug’urta* (insurance) and *Investitsiya* (investment)—drop to 2.67, driven by completeness failures rather than accuracy: the accuracy sub-score remains ≥ 3.67 across all 23 domains. Overall, 15 of 23 domains (65%) meet the inclusion threshold (≥ 3.5).

Table 8. Per-domain LLM-as-Judge quality scores (overall, 1–5 scale). Domains grouped by proximity to the training distribution.

Group	Domain	N	Comp.	Acc.	Sent.	Rel.	Overall
In-domain	Restoran	50	3.68	4.76	4.56	4.66	4.24
Near	Sayohat/Turizm	5	4.00	4.80	4.60	4.40	4.40
	Oziq-ovqat	15	3.47	4.33	4.47	4.13	3.80
Out-of-domain	Bank/Moliya	30	3.40	4.47	4.17	3.90	3.70
	Telekom.	25	3.04	4.36	3.64	3.64	3.32
	Ta’lim	20	2.90	4.20	3.80	3.45	3.25
Distant	Sug’urta	3	2.33	3.67	2.67	3.33	2.67
	Investitsiya	3	2.33	4.00	3.67	3.00	2.67

Figure 4 visualises the full 23-domain quality spectrum against the inclusion threshold, making the in-domain-to-distant gradient and the position of each domain relative to the threshold directly legible.

Annotation Pipeline Output. Batch annotation of the full 5,038-review corpus produces 9,884 aspects (average 1.96 per review) with a 100% JSON parse rate across all 23 domains. The polarity distribution is 75.1% positive, 16.5% negative, and 8.4% neutral—reflecting a natural skew in user-generated reviews. The aspect category distribution reveals a notable pattern: *boshqalar* (miscellaneous) dominates at 45.5%, followed by *ovqat* (17.9%), *xizmat* (15.9%), *muhit* (10.7%), and *narx* (9.0%).

5.5. Human Validation of Annotation Quality

Because both the silver annotations and their quality scores originate from models, we validate a subsample against native-speaker judgement. Two native Uzbek speakers—the first author and an independent trained annotator (acknowledged below)—annotated a stratified subsample of 150 reviews drawn from the 307 judged reviews (all 23 domains represented, deliberately oversampling distant and low-scoring domains), performing two tasks: (i) rating the model’s predicted annotations on the same five-dimension 1–5 rubric used by the LLM judge, and (ii) re-annotating a subset of 80 reviews from scratch to obtain human gold references.

Inter-annotator agreement. Rubric ratings were produced after a joint calibration pass on the 1–5 scale. Table 9 reports agreement between the two raters on all 150 reviews, measured with quadratic-weighted Cohen’s κ (which penalises large rating gaps more than adjacent ones) and Krippendorff’s α on the interval scale. Agreement is high and consistent across dimensions ($\kappa = 0.911$ – 0.970 , $\alpha = 0.911$ – 0.970), highest on sentiment correctness ($\kappa = 0.970$) and lowest on relevance ($\kappa = 0.911$); Spearman correlations between raters fall in the same band ($\rho = 0.920$ – 0.941). The two raters therefore apply the rubric consistently enough for their mean rating to serve as the reference against which the LLM judge is calibrated below.

Gold annotation followed a two-stage protocol. In the first, fully independent pass, the two annotators’ span-level agreement was low (exact-match F1 0.277, partial 0.415)—not because they disagreed about opinions, but because their segmentation and normalization conventions diverged (lemmatized vs. surface forms; naming the referent rather than the mention); polarity agreement on matched terms was perfect. In the second stage the annotators reconciled conventions and produced a consolidated reference with high residual span agreement (exact F1 0.927, polarity 100%), released as the reconciled double-annotated subset (*gold80*) alongside the silver corpus; because the reconciliation involved the same two annotators, we refrain from calling it fully adjudicated gold until an independent




Figure 4. Mean LLM-as-Judge overall quality score for all 23 business domains (review counts in parentheses), sorted by score. Blue bars meet the inclusion threshold of 3.5 (dashed line); gray bars fall below it. The gradient runs from in-domain (*Restoran/Ovqatlanish*) and vocabulary-adjacent domains at the top to specialised financial domains (*Sug'urta*, *Investitsiya*) at the bottom.

third annotator has resolved the residual disagreements. This two-stage outcome is itself informative: even for native speakers, *what counts as the aspect span* is the dominant source of ABSA annotation disagreement in Uzbek, while polarity is essentially unambiguous.

Judge calibration. Table 10 compares GPT-4o-mini's scores with the mean human rating on the same 150 reviews. Rank correlations are moderate to strong on the dimensions that drive dataset assembly—sentiment ($\rho = 0.81$), relevance ($\rho = 0.67$), and the tier-defining overall score ($\rho = 0.66$; all $p < 10^{-6}$)—and the judge's mean overall score (3.75) closely matches the human mean (3.84). Correlations are weaker for completeness and accuracy ($\rho \approx 0.43$ – 0.46), where the judge is systematically harsher on completeness (mean 3.30 vs. 4.04). We conclude that the LLM judge is a usable, though imperfect, proxy for native-speaker quality judgement, and that its quality tiers rank annotations in an order humans broadly agree with.

Table 9. Inter-annotator agreement between the two native-speaker raters on the 150-review rubric sample, per dimension. κ : quadratic-weighted Cohen’s kappa; α : Krippendorff’s alpha (interval); ρ : Spearman rank correlation.

Dimension	Weighted κ	Krippendorff’s α	Spearman ρ
Sentiment	0.970	0.970	0.941
Accuracy	0.936	0.936	0.920
Completeness	0.930	0.930	0.937
Overall	0.923	0.922	0.923
Relevance	0.911	0.911	0.925

Table 10. Calibration of the GPT-4o-mini judge against the mean rating of two native-speaker annotators on the same 150 reviews (Spearman ρ , all $p < 10^{-6}$; MAE on the 1–5 scale).

Dimension	ρ	MAE	Human mean	Judge mean
Sentiment	0.806	0.34	4.25	4.24
Relevance	0.667	0.50	3.85	3.97
Overall	0.657	0.55	3.84	3.75
Accuracy	0.433	0.54	4.19	4.44
Completeness	0.462	0.94	4.04	3.30

Model vs. human gold. Against the consolidated gold references, the model achieves exact-match ATE F1 of 0.209–0.224 and partial-match F1 of 0.425–0.433 (per annotator), with 100% polarity agreement on exactly matched terms. Two reference points contextualize these scores. First, the model’s partial-match agreement (≈ 0.43) matches the level two humans reach *before* convention reconciliation (0.415), while falling well short of the reconciled ceiling (0.93)—the gap is concentrated in span segmentation, not in opinion detection or polarity. Second, across domain-proximity groups the human-measured quality reproduces the judge’s gradient, falling from 0.244 exact pair F1 on general out-of-domain reviews to 0.170 on the most distant domains. We note the scope of this study: one of the two annotators is the first author, and the sample sizes per domain are small; a larger validation with fully external annotators is left for future work.

6. Discussion

The Completeness Bottleneck. Completeness (mean 3.32/5) is the weakest quality dimension: the model reliably extracts aspects present in the text but often misses domain-specific or subtle opinions, especially outside the restaurant domain. This is primarily due to the training data’s narrow coverage—only five aspect categories—forcing the model to map unfamiliar aspects to the generic *boshqalar* (miscellaneous) label, which dominates at 45.5%. Addressing this bottleneck will require domain-adaptive fine-tuning or expanding the aspect taxonomy to better capture out-of-domain phenomena.

Checkpoint Effects. R1 reasoning distillation conferred no advantage on structured ABSA in this setting: despite sharing the Qwen-7B backbone, the distilled checkpoint underperformed both Qwen and Llama. The zero-shot controls isolate part of the mechanism directly: unadapted, the R1 checkpoint’s reasoning traces break JSON parsing on more than half the evaluation examples (Table 5), so reasoning-oriented output habits demonstrably sit awkwardly against strict formatting—and even after fine-tuning suppresses the preamble, no residual benefit appears. Llama achieves the lowest loss and best sentiment accuracy, but Qwen is superior for reliable extraction and instruction following (100% parse rate, highest ATE recall). For annotation pipelines, robust JSON compliance and extraction coverage are more critical than marginal gains in sentiment metrics.

When is a 7B LLM Worth It? The encoder baselines sharpen the practical picture. A 110M monolingual encoder, fine-tuned in ten minutes on a laptop, shows no statistically significant difference from the 7–8B LLMs at aspect-term extraction—a reminder that for span identification in a morphologically rich language, targeted monolingual pre-training (BERTbek, TahrirchiBERT) remains remarkably effective. The LLMs justify their cost where the task is genuinely joint and generative: sentiment macro-F1 (+5–7 points over the encoders), unified term- and category-level output in a single generation, and the ability to follow natural-language instructions on unseen domains, which is what the multi-domain annotation pipeline exploits. Teams building Uzbek ABSA systems under tight compute budgets should therefore consider encoders for extraction-only workloads and reserve LLMs for joint prediction and annotation-generation scenarios.

Cost of the Pipeline. Because the practical case for model-assisted annotation rests on its economics, we report them explicitly. All QLoRA fine-tuning ran on a single 48 GB NVIDIA RTX A6000: 1,000 optimizer steps took 35 min (Qwen), 48 min (Llama), and 84 min (DeepSeek) of pure training time, or 65–107 min of wall clock including periodic evaluation. Annotating the 5,038-review multi-domain corpus with the fine-tuned Qwen model (Layer 1) ran at a median 4.3 s per review—roughly 840 reviews per hour—for approximately 6 GPU-hours end to end. Layer 2 judging is cheaper still: the 307-review stratified sample cost approximately \$0.05 in gpt-4o-mini API charges, which extrapolates to under \$1 to judge the corpus in full. The dominant remaining cost is therefore human: the validation study in Section 5.5 required two native speakers to rate 150 reviews and to annotate 80 from scratch, and it is this step—not GPU time or API spend—that bounds how much of a corpus can be verified. That asymmetry is the practical argument for the pipeline: machine annotation and machine judging are inexpensive enough to apply exhaustively, so scarce native-speaker effort can be spent on calibrating and spot-checking the automatic layers rather than on labeling every review.

Loss as a Model Selection Proxy. Validation loss ranking diverges from ABSA performance: Llama’s lowest loss does not translate to the best extraction, while Qwen outperforms on ATE F1 despite higher loss. Cross-entropy loss captures token-level likelihood but not structured-output quality such as JSON validity or aspect-boundary precision. This divergence is expected for structured-generation tasks, yet it is easy to overlook in practice; model selection for ABSA is therefore better guided by task-specific metrics than by validation loss alone.

Limitations. This study trains exclusively on restaurant reviews, without domain-adaptive fine-tuning; cross-domain performance should therefore be read as a demonstration of pipeline feasibility rather than a claim of uniform quality, and indeed we observe measurable degradation on distant domains. While we validate a stratified subset of the multi-domain corpus against native-speaker annotations (Section 5.5), the released dataset remains predominantly silver-standard: the human-verified portion is an 80-review subset produced by two native-speaker annotators under a two-stage reconciliation protocol, and the remaining annotations carry only model- and judge-derived quality estimates. Positive polarity dominates both datasets, and compute constraints meant the model comparison reuses the fixed evaluation split rather than a separately retrained held-out test set; for the same reason, statistical significance testing covers the encoder pair, while the LLM–encoder gaps are reported descriptively.

Several evaluation-design caveats remain. First, the 609-example split functions as a development set: evaluation loss on it is monitored during training (every 100 steps), and the same partition is then reused for the headline numbers, so the reported scores are optimistic relative to a genuinely untouched test set; a three-way train/development/test partition is the main methodological upgrade we would make with a larger corpus. Second,

although the split is deduplicated by review text (Section 5.3) and now controlled against zero-shot baselines (Section 5.2), robustness to random seeds is only partially characterized: a single replicate (for Qwen) bounds seed variation at roughly ± 0.01 F1, but multi-seed replication of every system—which Table 6 suggests matters at exactly the scale of the observed inter-model gaps—remains undone. Third, few-shot prompting sits unexplored between the zero-shot controls and full fine-tuning, so we cannot say how much of the QLoRA gain cheaper in-context adaptation could recover. Finally, the multi-domain case study (Section 5.4) was produced with the annotation engine fine-tuned under the earlier revision of the preprocessing pipeline; since its output quality is assessed directly—by the LLM judge and the native-speaker validation study—the case-study conclusions are unaffected, but the released silver annotations do not benefit from the improvements in Table 6, and regenerating them with the corrected model is planned for the next dataset version. Tokenizer efficiency and larger-scale, convention-reconciled human verification remain for future work.

7. Conclusions

This paper presented, to our knowledge, the first systematic study of parameter-efficient fine-tuning of open-weight LLMs for Uzbek ABSA. We showed that QLoRA enables effective joint aspect extraction and sentiment classification via structured JSON generation—at least doubling extraction F1 and tripling end-to-end pair F1 over unadapted zero-shot controls—and that the fine-tuned LLMs are competitive with fine-tuned Uzbek BERT encoders. Our comparison revealed a clear architectural trade-off: Qwen 2.5-7B excels at reliable aspect extraction and instruction following, whereas Llama 3.1-8B achieves the highest precision for end-to-end sentiment matching, with the two effectively tied on end-to-end F1. We also found that validation loss is an unreliable proxy for downstream structured-extraction performance. Finally, we showed that a model trained only on restaurant reviews can produce usable annotations on more distant domains, though quality degrades measurably with domain distance. Our three-layer pipeline, together with a human-validation study of its output, produced a quality-scored multi-domain silver-standard resource with a gold-verified subset, offering a practical route to scalable low-resource dataset creation. Future work will focus on expanding the aspect taxonomy, on domain-adaptive fine-tuning to improve extraction completeness across diverse business domains, and on injecting the morphosyntactic structure now available from Uzbek Universal Dependencies treebanks and parsers [31–33] into the extraction step, since span segmentation—not opinion detection—is where both our models and our human annotators disagreed most (Section 5.5).

Author Contributions: Conceptualization, S.M. and M.A.; methodology, S.M. and M.A.; software, S.M.; validation, S.M. and G.D.; formal analysis, S.M. and G.D.; investigation, S.M. and G.D.; resources, S.M.; data curation, S.M.; writing—original draft preparation, S.M.; writing—review and editing, S.M., G.D. and M.A.; visualization, S.M.; supervision, M.A.; project administration, S.M.; funding acquisition, not applicable. All authors have read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Institutional Review Board Statement: Not applicable.

Informed Consent Statement: Not applicable. The study analyses publicly accessible user-generated business reviews and does not involve direct interaction with human participants. The reviews were used with explicit permission from the platform owner (see Section 7).

Data Availability Statement: The primary Uzbek aspect-based sentiment analysis dataset used for model fine-tuning is publicly available on Hugging Face at <https://huggingface.co/datasets/>

[Sanatbek/aspect-based-sentiment-analysis-uzbek](#). The multi-domain Uzbek business review corpus generated in this study, including the 80-review reconciled double-annotated subset (Section 5.5), is publicly available on Zenodo at <https://zenodo.org/records/22146677> with DOI: [10.5281/zenodo.22146677](https://doi.org/10.5281/zenodo.22146677). The released records contain the review text, business category and name, the star rating, and the annotation and quality metadata; contributor user names, review URLs, and posting timestamps present in the source platform are deliberately excluded. Code for the encoder baselines, the human-validation analysis, and the significance tests is available in the project repository.

Acknowledgments: We extend our heartfelt gratitude to Sardor Berdiyev, CEO and Co-founder of Commeta, for graciously granting permission to use publicly available reviews from the sharh.commeta.uz platform exclusively for academic research purposes. We thank Jaloliddin Rajabov for serving as the second native-speaker annotator in the human validation study (Section 5.5). We also acknowledge the computational resources provided by the National University of Uzbekistan. During the preparation of this manuscript/study, the authors used Anthropic Claude (Sonnet 5 and Opus 5) for the purposes of paraphrasing, language editing, and manuscript revision. The authors have reviewed and edited the output and take full responsibility for the content of this publication.

Conflicts of Interest: The authors declare no conflicts of interest.

Abbreviations

The following abbreviations are used in this manuscript:

ABSA	Aspect-Based Sentiment Analysis
ACD	Aspect Category Detection
ASC	Aspect Sentiment Classification
ATE	Aspect Term Extraction
QLoRA	Quantized Low-Rank Adaptation
RQ	Research Question

Appendix A. Full Per-Domain LLM-as-Judge Quality Scores

Table A1 reports LLM-as-Judge quality scores for all 23 business domains in the stratified evaluation sample (307 reviews total). The main text (Table 8) presents a selected subset grouped by training-distribution proximity; this table provides the complete picture, sorted by overall quality score in descending order. Scores are on a 1–5 Likert scale across five dimensions: Completeness, Accuracy (hallucination control), Sentiment Correctness, Relevance (category–domain fit), and Overall. Domains with $N \leq 5$ should be interpreted with caution due to small sample sizes. The “Include” threshold (Overall ≥ 3.5) is shown as a dashed reference.

Table A1. LLM-as-Judge (GPT-4o-mini) quality scores for all 23 business domains, sorted by Overall score (descending). Stratified evaluation sample: 307 reviews. Scale: 1–5 (higher is better). * Small sample ($N \leq 5$): interpret with caution.

Domain (Uzbek / English)	N	Comp.	Acc.	Sent.	Rel.	Overall
Bozor/BSC (Market/Shopping centre)*	3	3.67	5.00	5.00	5.00	4.67
Sayohat/Turizm (Tourism/Travel)*	5	4.00	4.80	4.60	4.40	4.40
Din/Madaniyat (Religion/Culture)*	3	4.33	5.00	4.00	5.00	4.33
Yetkazib berish (Delivery)*	3	3.67	5.00	4.67	4.33	4.33
Restoran/Ovqatlanish (Restaurant/Dining)	50	3.68	4.76	4.56	4.66	4.24
Go'zallik (Beauty services)*	3	3.33	4.67	5.00	4.00	4.00
Texnologiya/Media (Technology/Media)	8	2.88	4.50	5.00	3.88	3.88
Boshqa (Other/Miscellaneous)	37	3.43	4.54	3.89	4.38	3.81
Elektron tijorat (E-commerce)	21	3.43	4.48	4.10	4.14	3.81
Oziq-ovqat do'konlari (Grocery stores)	15	3.47	4.33	4.47	4.13	3.80
Tibbiyot/Sog'liqni saqlash (Healthcare)	25	3.40	4.60	4.20	4.36	3.80
Bank/Moliya (Banking/Finance)	30	3.40	4.47	4.17	3.90	3.70
Ko'ngilochar (Entertainment)*	3	3.33	4.67	4.33	4.00	3.67
Transport/Yo'l (Transport/Roads)*	3	3.33	4.00	4.00	4.00	3.67
Gul/Sovg'a (Flowers/Gifts)	12	3.08	4.08	4.83	3.67	3.58
To'lov tizimlari (Payment systems)	24	3.00	4.29	4.38	3.54	3.54
<i>Below inclusion threshold (Overall < 3.5)</i>						
Davlat xizmatlari (Government services)*	3	3.00	4.00	4.33	3.67	3.33
Telekommunikatsiya (Telecommunications)	25	3.04	4.36	3.64	3.64	3.32
Ta'lim (Education)	20	2.90	4.20	3.80	3.45	3.25
Sport/Fitnes (Gyms/Fitness)*	4	2.75	4.25	4.25	3.00	3.25
Kitob/Nashriyot (Books/Publishing)*	4	3.00	4.25	4.00	3.00	3.00
Sug'urta (Insurance)*	3	2.33	3.67	2.67	3.33	2.67
Investitsiya/Trading (Investment)*	3	2.33	4.00	3.67	3.00	2.67

Appendix B. Bootstrap Confidence Intervals for All Systems

Table A2 reports exact percentile-bootstrap 95% confidence intervals ($B=10,000$ re-samples over the 609 evaluation reviews) for every score in Table 4. Half-widths range from ± 0.023 (Llama sentiment accuracy) to ± 0.037 (DeepSeek pair F1), and the intervals of the top-ranked systems overlap substantially on all four metrics—the reason the paired-bootstrap comparisons in Section 5.1 detect so few significant differences. Note that these are per-system intervals on absolute scores; because the systems are evaluated on identical reviews, the *paired* tests in Section 5.1 are more powerful than a visual comparison of these intervals would suggest.

Table A2. Percentile-bootstrap 95% confidence intervals ($B=10,000$, resampling reviews) for each system on the 609-example evaluation set. Estimates match Table 4.

System	ATE exact F1	ATE partial F1	Pair F1	Sent. Acc.
<i>QLoRA fine-tuned LLMs</i>				
Qwen 2.5-7B	0.708 [0.674, 0.740]	0.802 [0.775, 0.829]	0.645 [0.608, 0.680]	0.911 [0.886, 0.935]
Llama 3.1-8B	0.704 [0.670, 0.737]	0.801 [0.773, 0.828]	0.651 [0.615, 0.687]	0.925 [0.901, 0.947]
DeepSeek-R1-7B	0.678 [0.642, 0.714]	0.787 [0.758, 0.817]	0.603 [0.566, 0.640]	0.891 [0.860, 0.918]
<i>Fine-tuned Uzbek encoders</i>				
BERTbek	0.692 [0.657, 0.726]	0.800 [0.774, 0.826]	0.608 [0.571, 0.644]	0.879 [0.851, 0.906]
TahrirchiBERT	0.701 [0.667, 0.734]	0.787 [0.758, 0.814]	0.623 [0.586, 0.657]	0.889 [0.862, 0.915]

References

- Pontiki, M.; Galanis, D.; Pavlopoulos, J.; Papageorgiou, H.; Androutsopoulos, I.; Manandhar, S. SemEval-2014 Task 4: Aspect Based Sentiment Analysis. In Proceedings of the Proceedings of the 8th International Workshop on Semantic Evaluation (SemEval 2014); Nakov, P.; Zesch, T., Eds., Dublin, Ireland, August 2014; pp. 27–35. <https://doi.org/10.3115/v1/S14-2004>.

2. Qwen.; Yang, A.; Yang, B.; Zhang, B.; Hui, B.; Zheng, B.; Yu, B.; Li, C.; Liu, D.; Huang, F.; et al. Qwen2.5 Technical Report. *arXiv preprint* **2024**, [arXiv:cs.CL/2412.15115]. 799
3. Grattafiori, A.; Dubey, A.; Jauhri, A.; et al. The Llama 3 herd of models. *arXiv preprint* **2024**, [arXiv:cs.AI/2407.21783]. 800
4. Guo, D.; Yang, D.; Zhang, H.; Song, J.; Wang, P.; Zhu, Q.; Xu, R.; Zhang, R.; Ma, S.; Bi, X.; et al. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature* **2025**, *645*, 633–638. <https://doi.org/10.1038/s41586-025-09422-z>. 801
5. Dettmers, T.; Pagnoni, A.; Holtzman, A.; Zettlemoyer, L. QLORA: efficient finetuning of quantized LLMs. In Proceedings of the Proceedings of the 37th International Conference on Neural Information Processing Systems, Red Hook, NY, USA, 2023; NIPS '23. 802
6. Ma, D.; Li, S.; Wu, F.; Xie, X.; Wang, H. Exploring Sequence-to-Sequence Learning in Aspect Term Extraction. In Proceedings of the Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics; Korhonen, A.; Traum, D.; Màrquez, L., Eds., Florence, Italy, July 2019; pp. 3538–3547. <https://doi.org/10.18653/v1/P19-1344>. 803
7. Hu, M.; Zhao, S.; Guo, H.; Xue, C.; Gao, H.; Gao, T.; Cheng, R.; Su, Z. Multi-Label Few-Shot Learning for Aspect Category Detection. In Proceedings of the Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers); Zong, C.; Xia, F.; Li, W.; Navigli, R., Eds., Online, August 2021; pp. 6330–6340. <https://doi.org/10.18653/v1/2021.acl-long.495>. 804
8. Sun, X.; Mi, Y.; Li, H. Enhancing Aspect Sentiment Classification with Dual-Channel Graph Convolutional Network. *ACM Trans. Intell. Syst. Technol.* **2025**, *16*. <https://doi.org/10.1145/3721844>. 805
9. Hua, Y.C.; Denny, P.; Wicker, J.; Taskova, K. A systematic review of aspect-based sentiment analysis: domains, methods, and trends. *Artificial Intelligence Review* **2024**, *57*, 296. <https://doi.org/10.1007/s10462-024-10906-z>. 806
10. Zohreh Madhoushi, Z.M.; Hamdan, A.R.; Zainudin, S. Aspect-Based Sentiment Analysis Methods in Recent Years. *Asia-Pacific Journal of Information Technology; Multimedia* **2019**, *08*, 79–96. <https://doi.org/10.17576/apjitm-2019-0801-07>. 807
11. Chouikhi, H.; Alsuhaibani, M.; Jarray, F. BERT-Based Joint Model for Aspect Term Extraction and Aspect Polarity Detection in Arabic Text. *Electronics* **2023**, *12*. <https://doi.org/10.3390/electronics12030515>. 808
12. Zou, F.; Zou, L.; Guo, F.; Wang, X.; Weng, J.; Fang, T.; Jiang, H.; Wu, X. Enhanced Quantum-Inspired Deep Learning with Multi-Head Attention and Contrastive Learning for Text-Based Dialogue Sentiment Classification. *Big Data and Cognitive Computing* **2026**, *10*. <https://doi.org/10.3390/bdcc10050161>. 809
13. Liao, S.w.; Wang, C.S.; Yeh, C.C.; Lin, J.W. Aspect-Based Sentiment Analysis with Enhanced Opinion Tree Parsing and Parameter-Efficient Fine-Tuning for Edge AI. *Electronics* **2025**, *14*. <https://doi.org/10.3390/electronics14040690>. 810
14. Alotaibi, A.; Nadeem, F. An Unsupervised Integrated Framework for Arabic Aspect-Based Sentiment Analysis and Abstractive Text Summarization of Traffic Services Using Transformer Models. *Smart Cities* **2025**, *8*. <https://doi.org/10.3390/smartcities8020062>. 811
15. Scaria, K.; Gupta, H.; Goyal, S.; Sawant, S.; Mishra, S.; Baral, C. InstructABSA: Instruction Learning for Aspect Based Sentiment Analysis. In Proceedings of the Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers); Duh, K.; Gomez, H.; Bethard, S., Eds., Mexico City, Mexico, June 2024; pp. 720–736. <https://doi.org/10.18653/v1/2024.naacl-short.63>. 812
16. Simmering, P.F.; Huoviala, P. Large Language Models for Aspect-Based Sentiment Analysis. *arXiv preprint arXiv:2310.18025* **2023**. <https://doi.org/10.48550/arXiv.2310.18025>. 813
17. Šmíd, J.; Priban, P.; Kral, P. LLaMA-Based Models for Aspect-Based Sentiment Analysis. In Proceedings of the Proceedings of the 14th Workshop on Computational Approaches to Subjectivity, Sentiment, & Social Media Analysis; De Clercq, O.; Barriere, V.; Barnes, J.; Klinger, R.; Sedoc, J.; Tafreshi, S., Eds., Bangkok, Thailand, August 2024; pp. 63–70. <https://doi.org/10.18653/v1/2024.wassa-1.6>. 814
18. Wu, C.; Ma, B.; Zhang, Z.; Deng, N.; He, Y.; Xue, Y. Evaluating zero-shot multilingual aspect-based sentiment analysis with large language models. *International Journal of Machine Learning and Cybernetics* **2025**, *16*, 8079–8101. <https://doi.org/10.1007/s13042-025-02711-z>. 815
19. Liskowski, P.; Jankowski, K. Large-Scale Aspect-Based Sentiment Analysis with Reasoning-Infused LLMs. *arXiv preprint arXiv:2601.03940* **2026**. <https://doi.org/10.48550/arXiv.2601.03940>. 816
20. Touvron, H.; Lavril, T.; Izacard, G.; Martinet, X.; Lachaux, M.A.; Lacroix, T.; Rozière, B.; Goyal, N.; Hambro, E.; Azhar, F.; et al. LLaMA: Open and efficient foundation language models. *arXiv preprint* **2023**, [arXiv:cs.CL/2302.13971]. 817
21. Hu, E.J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; Chen, W. LoRA: Low-Rank Adaptation of large language models. *arXiv preprint arXiv:2106.09685* **2021**, [arXiv:cs.CL/2106.09685]. 818
22. Zhang, K.; Lin, X.; Wang, M.; Yu, H.; Chen, L. Named Entity Recognition Method for Natural Disaster Emergencies Based on Instruction Tuning and Graph Retrieval-Augmented Generation. *Big Data and Cognitive Computing* **2026**, *10*. <https://doi.org/10.3390/bdcc10060185>. 819
23. Pingua, B.; Sahoo, A.; Kandpal, M.; Murmu, D.; Rautaray, J.; Barik, R.K.; Saikia, M.J. Medical LLMs: Fine-Tuning vs. Retrieval-Augmented Generation. *Bioengineering* **2025**, *12*. <https://doi.org/10.3390/bioengineering12070687>. 820

24. Abdellaoui, I.; Ibrahimi, A.; El Bouni, M.A.; Mourhir, A.; Driouech, S.; Aghzal, M. Investigating Offensive Language Detection in a Low-Resource Setting with a Robustness Perspective. *Big Data and Cognitive Computing* **2024**, *8*. <https://doi.org/10.3390/bdcc8120170>.
25. Matlatipov, G.; Vetulani, Z., Representation of Uzbek Morphology in Prolog. In *Aspects of Natural Language Processing: Essays Dedicated to Leonard Bolc on the Occasion of His 75th Birthday*; Springer Berlin Heidelberg: Berlin, Heidelberg, 2009; pp. 83–110. https://doi.org/10.1007/978-3-642-04735-0_4.
26. Matlatipov, S.; Tukeyev, U.; Aripov, M. Towards the Uzbek Language Endings as a Language Resource. In *Proceedings of the Advances in Computational Collective Intelligence*; Hernes, M.; Wojtkiewicz, K.; Szczerbicki, E., Eds., Cham, 2020; pp. 729–740.
27. Aripov, M.; Khakimov, M.; Matlatipov, S.; Sirojiddinov, Z. Analysis and Processing of the Uzbek Language on the Multi-language Modelled Computer Translator Technology. In *Proceedings of the Human Language Technology. Challenges for Computer Science and Linguistics*; Vetulani, Z.; Paroubek, P.; Kubis, M., Eds., Cham, 2022; pp. 81–95.
28. Sharipov, M.; Kuriyozov, E.; Yuldashev, O.; Sobirov, O. UzbekTagger: The rule-based POS tagger for Uzbek language. *arXiv preprint arXiv:2301.12711* **2023**, [arXiv:cs.CL/2301.12711].
29. Salaev, U. UzMorphAnalyser: A morphological analysis model for the Uzbek language using inflectional endings. *AIP Conference Proceedings* **2024**, 3244, 030058. <https://doi.org/10.1063/5.0241461>.
30. Salaev, U.; Kuriyozov, E.; Gómez-Rodríguez, C. A machine transliteration tool between Uzbek alphabets. In *Proceedings of the Proceedings of the International Conference and Workshop on Agglutinative Language Technologies as a Challenge of Natural Language Processing (ALT/NLP)*, Koper, Slovenia, 2022. <https://doi.org/10.48550/arXiv.2205.09578>.
31. Matlatipov, S.; Aripov, M.; Bobokandov, M.; Matlatipov, G. Annotated universal dependencies dataset for literary and educational Uzbek texts. *Data in Brief* **2026**, *66*, 112857. <https://doi.org/https://doi.org/10.1016/j.dib.2026.112857>.
32. Matlatipov, S.; Ismoilova, M.; Aripov, M. Graph-Based vs Transition-Based Dependency Parsing for Uzbek. In *Proceedings of the 2026 IEEE 27th International Conference of Young Professionals in Electron Devices and Materials (EDM)*, 2026, pp. 1–6. <https://doi.org/10.1109/EDM69524.2026.11632281>.
33. Matlatipov, S. Towards Robust Uzbek Neural Dependency Parsing: Cross-Treebank Training. In *Proceedings of the Natural Language Processing and Information Systems*; Cabrio, E.; Monteiro, E., Eds., Cham, 2027; pp. 204–218.
34. Yusufu, A.; Liu, J.; Aziz, K.; Ainiwaer, A.; Li, B.; Li, F.; Ji, D.; Yusufu, A. LASQ: A low-resource aspect-based sentiment quadruple extraction dataset. *Journal of King Saud University Computer and Information Sciences* **2026**, *38*, 337. <https://doi.org/10.1007/s44443-026-00701-x>.
35. Kuriyozov, E.; Matlatipov, S. Building a New Sentiment Analysis Dataset for Uzbek Language and Creating Baseline Models. *Proceedings* **2019**, 21. <https://doi.org/10.3390/proceedings2019021037>.
36. Kuriyozov, E.; Matlatipov, S.; Alonso, M.A.; Gómez-Rodríguez, C. Construction and Evaluation of Sentiment Datasets for Low-Resource Languages: The Case of Uzbek. *Lecture Notes in Artificial Intelligence* **2022**, 13212, 232–243. https://doi.org/10.1007/978-3-031-05328-3_15.
37. Matlatipov, S.; Rahimboeva, H.; Rajabov, J.; Kuriyozov, E. Uzbek Sentiment Analysis based on local Restaurant Reviews. *arXiv preprint arXiv:2205.15930* **2022**, [arXiv:cs.CL/2205.15930].
38. Matlatipov, S.G.; Rajabov, J.; Kuriyozov, E.; Aripov, M. UzABSA: Aspect-Based Sentiment Analysis for the Uzbek Language. In *Proceedings of the Proceedings of the 3rd Annual Meeting of the Special Interest Group on Under-resourced Languages @ LREC-COLING 2024*; Melero, M.; Sakti, S.; Soria, C., Eds., Torino, Italia, May 2024; pp. 394–403.
39. Kuriyozov, E.; Vilares, D.; Gómez-Rodríguez, C. BERTbek: A Pretrained Language Model for Uzbek. In *Proceedings of the Proceedings of the 3rd Annual Meeting of the Special Interest Group on Under-resourced Languages @ LREC-COLING 2024*; Melero, M.; Sakti, S.; Soria, C., Eds., Torino, Italia, May 2024; pp. 33–44.
40. Mamasaidov, M.; Shopulatov, A. TahrirchiBERT base: A RoBERTa-based Language Model Pre-trained on 5B Tokens of Uzbek Latin-script Text. Hugging Face model card, 2023. Available online: <https://huggingface.co/tahrirchi/tahrirchi-bert-base> (accessed on 8 July 2026).
41. Zheng, L.; Chiang, W.L.; Sheng, Y.; Zhuang, S.; Wu, Z.; Zhuang, Y.; Lin, Z.; Li, Z.; Li, D.; Xing, E.P.; et al. Judging LLM-as-a-judge with MT-bench and Chatbot Arena. In *Proceedings of the Proceedings of the 37th International Conference on Neural Information Processing Systems*, Red Hook, NY, USA, 2023; NIPS '23.
42. Li, Q.; Dou, S.; Shao, K.; Chen, C.; Hu, H. Evaluating Scoring Bias in LLM-as-a-Judge, 2025, [arXiv:2506.22316].
43. Hellwig, N.C.; Fehle, J.; Kruschwitz, U.; Wolff, C. LLM-as-an-Annotator: Training Lightweight Models with LLM-Annotated Examples for Aspect Sentiment Tuple Prediction. In *Proceedings of the Proceedings of the Fifteenth Language Resources and Evaluation Conference (LREC 2026)*; Piperidis, S.; Bel, N.; van den Heuvel, H.; Ide, N.; Krek, S.; Toral, A., Eds., Palma, Mallorca, Spain, May 2026; pp. 7955–7972. <https://doi.org/10.63317/43srcdyc52cd>.
44. Koehn, P. Statistical Significance Tests for Machine Translation Evaluation. In *Proceedings of the Proceedings of the 2004 Conference on Empirical Methods in Natural Language Processing*; Lin, D.; Wu, D., Eds., Barcelona, Spain, July 2004; pp. 388–395.

Disclaimer/Publisher's Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content.

907
908
909